

# Mobile Programming

Ian Batten

[igb@batten.eu.org](mailto:igb@batten.eu.org)

# Today's Content

- How did the exercise go?
  - A variety of things mean I haven't yet looked at the submissions in detail, but what I have seen is very good.
- Thread and Process Models (general)
- Thread and Process Models (Android)
- Services
- Broadcast

# Process Models

- Android is Linux, and although the userland is very different, the kernel and the basic structure of processes is the same.
- So it is worth our while walking through the Unix process model to see how Android applications fit in.

# Why do we care?

- We want to be able to (for example)
  - run computation while still responding to UI events
  - read from multiple sensors in arbitrary order
- In more complex cases we want to:
  - run services which distribute updates to multiple applications
  - share computations amongst multiple clients
- In general, we want to block on some “wait channel” while continuing to execute other code

# Unix processes

- Process 1 (historically “init”) is created ex nihilo by the kernel. The kernel knows what a process looks like, so can make one from scratch.
- All other processes are the result of calling `fork()` or one of its equivalents to copy the existing process.
- Processes have their own address space, their own file descriptors, their own signal handlers.
- Processes are also “expensive”: `fork()` is slow and significant resources are consumed even by quiescent processes.

# fork()

- `fork()` can be called at any time, by any process.
- Assuming there is no error, `fork()` returns **twice**:
  - It returns a “process identifier”, a small integer of type `pid_t`, into the calling (“parent”) process;
  - It returns 0 into a new (“child”) process which has the process id returned into the parent.
- Modulo some minor details, the processes are otherwise the same.

# fork() and exec()

- The usual thing that happens is that fork() is called, then some minor housekeeping is done to arrange the correct file descriptors on stdin (0), stdout (1) and stderr (2), and then a variant of exec() is called.
- exec() replaces the running process with a new binary, but leaves open file descriptors (and other process information) unchanged.

# In action

```
/* headers deleted */

int
main (int argc, char **argv) {
    int outfd = open ("/tmp/outputfile", O_WRONLY | O_CREAT | O_TRUNC, 0644);
    int pid = fork ();
    if (pid == 0) {
        ++argv;
        dup2 (outfd, 1);
        dup2 (outfd, 2);
        execvp (argv[0], argv);
    } else if (pid > 0) {
        int loc;
        wait (&loc);
        fprintf (stderr, "program %s terminated, 0x%x\n", argv[1], loc);
    }
}
```

```
ians-macbook-air:MUC13 igb$ date; ./exec date
Tue 18 Feb 2014 10:32:53 GMT
program date terminated, 0x0
ians-macbook-air:MUC13 igb$ cat /tmp/outputfile
Tue 18 Feb 2014 10:32:53 GMT
ians-macbook-air:MUC13 igb$
```



# This is the “old way” to...

- Write a daemon (open a socket, listen on that socket, when a new connection arrives fork and process the connection, in the parent close and got back to listening)
- Process in the background (fork, start the background task going, communicate via interprocess-communication)

# Interprocess Communications

- pipes: pseudo files that can be written in one process and read in another process, provided the processes are related
- named pipes (aka fifos) which can be found in the filesystem, although they are not files
- Unix domain sockets, which are bidirectional
- Message queues
- Using TCP (or other networking) in loopback
- etc

# Pipes

```
int
main (int argc, char **argv) {
    int pipefd[2];
    pipe (pipefd);
    int pid = fork ();
    if (pid == 0) {
        close (pipefd[0]);
        dup2 (pipefd[1], 1);
        dup2 (pipefd[1], 2);
        fprintf (stdout, "Hello to the other process\n");
        exit (0);
    } else if (pid > 0) {
        char buffer[BUFSIZ];
        char *cr;
        int bytes;
        close (pipefd[1]);
        while ((bytes = read (pipefd[0], buffer, sizeof (buffer))) > 0) {
            buffer[bytes - 1] = '\0';
            fprintf (stdout, "the other end sent me: \"%s\" (that was nice)\n", buffer);
        }
        int loc;
        wait (&loc);
        fprintf (stderr, "other end terminated, 0x%x\n", loc);
    }
}
```

```
ians-macbook-air:MUC13 igb$ ./pipe
the other end sent me: "Hello to the other process" (that was nice)
other end terminated, 0x0
ians-macbook-air:MUC13 igb$
```

# What you really want is two “things” working on the same memory

- Unix has explicit “shared memory”, which is a legacy of work done by AT&T in the late 1980s and is usually regarded as a bit messy.
- Two processes can obtain a “key” to use a region of memory as though it is their own. Shared memory has its own namespace.
- There are some locking facilities available, in the shape of semaphores, again in their own namespace.
- It’s also possible to do this with creative use of `mmap ( )`

# Hence threads

- Multiple **threads** run in
  - the same address space
  - over the same memory
  - but with their own stack and their own program counter
- So method-local variables are thread-local, while global (or class-local) variables are shared between all threads

# Java (general-style)

- A class is written which implements “Runnable” or extends “Thread”
- In each case, it has a “start” method which starts it running in parallel with the lead thread.
- “synchronized” methods or “synchronized (obj) {...}” statements provide mutual exclusion
- You get to write portable code, but quite what is does is not clearly defined

# Threads: Why?

- We said at the outset that we want threads so that we can block on some event while continuing to execute code
- But this isn't the only reason.

# Clock Speeds Versus Core Count

- There have for a long time been multiple-processor computers to get more power out of the same clockspeed.
  - VAX 11/782 is two VAX 11/780s in a box running VMS or various Unix derivatives
  - Honeywell DPS-8/70M could go up to six processors running Multics
  - Sequent Symmetry and its successors used up to 30 x86 processors running a hacked Unix
  - Sun Starfires had up to 64 (later 128) UltraSparc processors
- All of these are pretty exotic for their era



# The problem

- The problem is that power consumption and heat usually scale non-linearly with clockspeed, and therefore doubling the clockspeed is much harder than doubling the core-count on one die.
- A common trick is to manufacture dies with (say) 8 cores in them, and then sell them as 1-, 2-, 4-, 6- and 8-way processors depending on how many actually work.

# So...

- A typical machine, particularly one which needs to be power- and heat-efficient, will not have one fast processor.
- Rather, it will have multiple processors, each of them rather lower powered than the total claim the manufacturer makes.
- A single-threaded, compute-bound application will only be able to take advantage of one core.
- Conversely, an application which does only run on one core will permit other applications to run smoothly, and will (in isolation) reduce power consumption.

# Running One Thread

- If you just start() a thread, you have little control over how it executes. But on Android:
- The run() method is called within the Runnable
- The run() method runs on another thread, that is not the UI thread
- The thread is ideally deprioritised

```
android.os.Process.setThreadPriority(android.os.Process.THREAD_PRIORITY_BACKGROUND);
```

# Runnable does not run on the UI Thread

- A Runnable thread invoked via start() cannot directly manipulate the UI objects
- If you want to do this easily, you use an AsyncThread as already described to run one Runnable and have the magic for handing things to the UI done for you
- If you want to do this from a more complex thread object, you need to use Handlers and Messages, which we will discuss next week

# Running multiple threads

- The fastest way to perform a computation or data reduction is to run a thread on each core
  - Switching between threads wastes resources on that core
  - Running fewer threads “wastes” (at least, leaves idle) cores.
- Standard pattern is to have a pool of threads, and execute tasks from a queue of Runnableables

# ThreadPoolExecutor

- Simple interfaces:
  - `newCachedThreadPool()` – start new threads, but reclaim ones that have finished
  - `newFixedThreadPool(int)` – fixed size pool
  - `newSingleThreadExecutor()` – run one thread in the background, but with code that is easily changed to something else
- `ThreadPoolExecutor(...)` allows you to tune everything: pool size, queuing, reclamation, scheduling...

# What about processes?

- You can group threads into processes and control which process contains which thread, and then use IPC to communicate between them
- If you need to do this, the `android:process` element in the application manifest allows you to assign activities to distinct processes

# OK, so what might use this?

- Services – long-running activities which are not necessarily part of just one application, and are not killed when they aren't on screen
- Services can be started so that they run indefinitely, potentially in a process of their own
- Services can be bound so that they run for as long as the service is being used, and then are reclaimed.
- An obvious example would be a location service of some sort.



# Intents

- They are instructions to perform an action, or that something that happened which might trigger an action
- An Intent is an action, plus an action-specific URI as additional data (a number to call, a contact to display).
- BroadcastReceivers receive intents that have been broadcast, and therefore filter which ones they get.
- Services have intents unicast to them.

# Service methods

- Services subclass Service or IntentService,
- Service overrides:
  - onStartCommand() for when someone called startService(Intent)
  - onBind() for when someone calls BindService(Intent)
  - onCreate() to be called on startup
  - onDestroy() for when the service shuts down

# Problems with Service

- The service runs on the main thread, so needs to create a new thread for each start request (ie, each application that starts the service) unless you know it's going to complete quickly
- This is like a Unix daemon, if you have written one.
- Without care, it can spawn a lot of threads.

# IntentService

- IntentService creates one thread, and queues start requests up. It restricts the resources that can be used.
- Just has to override OnHandleIntent

# IntentService sample

```
public class HelloIntentService extends IntentService {

    /**
     * A constructor is required, and must call the super IntentService(String)
     * constructor with a name for the worker thread.
     */
    public HelloIntentService() {
        super("HelloIntentService");
    }

    /**
     * The IntentService calls this method from the default worker thread with
     * the intent that started the service. When this method returns, IntentService
     * stops the service, as appropriate.
     */
    @Override
    protected void onHandleIntent(Intent intent) {
        // Normally we would do some work here, like download a file.
        // For our sample, we just sleep for 5 seconds.
        long endTime = System.currentTimeMillis() + 5*1000;
        while (System.currentTimeMillis() < endTime) {
            synchronized (this) {
                try {
                    wait(endTime - System.currentTimeMillis());
                } catch (Exception e) {
                }
            }
        }
    }
}
```

# Bound Service Example

```
public class LocalService extends Service {
    // Binder given to clients
    private final IBinder mBinder = new LocalBinder();
    // Random number generator
    private final Random mGenerator = new Random();

    /**
     * Class used for the client Binder. Because we know this service always
     * runs in the same process as its clients, we don't need to deal with IPC.
     */
    public class LocalBinder extends Binder {
        LocalService getService() {
            // Return this instance of LocalService so clients can call public
            methods
            return LocalService.this;
        }
    }

    @Override
    public IBinder onBind(Intent intent) {
        return mBinder;
    }

    /** method for clients */
    public int getRandomNumber() {
        return mGenerator.nextInt(100);
    }
}
```

# BroadcastReceivers

- This all assumes that we are going to run services on request from clients
- Another common pattern is that a service performs some task, and then tells some number of other applications (which the service may not know by name) that something has happened.
- Example: geofencing
- Example: Battery Low condition

# System Intents

```
<application
    android:icon="@drawable/ic_launcher"
    android:label="@string/app_name"
    android:theme="@style/AppTheme" >

    <receiver android:name="MyReceiver">
        <intent-filter>
            <action android:name="android.intent.action.BATTERY_LOW">
        </action>
        </intent-filter>
    </receiver>

</application>

public class MyReceiver extends BroadcastReceiver {

    @Override
    public void onReceive(Context context, Intent intent) {
        Toast.makeText(context, "Intent Detected.", Toast.LENGTH_LONG).show();
    }
}
```



# Custom Intents

```
public void broadcastIntent(View view)
{
    Intent intent = new Intent();
    intent.setAction("com.tutorialspoint.CUSTOM_INTENT");
    sendBroadcast(intent);
}
```

```
<application
    android:icon="@drawable/ic_launcher"
    android:label="@string/app_name"
    android:theme="@style/AppTheme" >

    <receiver android:name="MyReceiver">
        <intent-filter>
            <action android:name="com.tutorialspoint.CUSTOM_INTENT">
            </action>
        </intent-filter>
    </receiver>

</application>
```

# Intents can carry data

- Constructor can include a URI, which you build according to your requirements
- We will come back to intents next week, when we talk about using them to chain applications together.