

Buffer overflows

Secure Programming
Lecture 6

Where are we?

- We have started looking at common classes of vulnerabilities
- Today, we begin looking at memory corruption vulnerabilities, focusing on stack-based overflows

Buffer overflow

- Program writes data to a buffer and overruns the buffer's boundary overwriting adjacent memory
- Mostly relevant in C and C++ programs
- Memory-safe languages are not affected
 - Dynamic bound checks (e.g., Java)
 - Automatic resizing (e.g., Perl)

Usual suspects

Buffer overflows are associated with a number of insecure library functions that make it easy to overwrite buffers

String manipulation functions without boundary checking

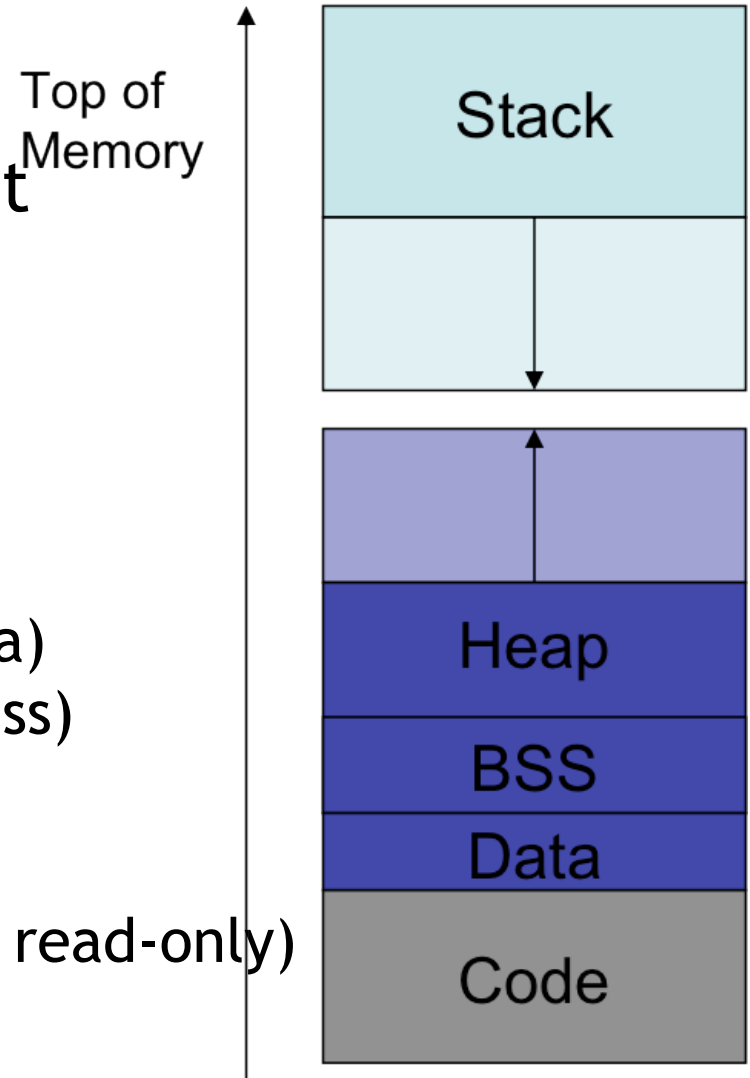
- `gets`
- `strcat`
- `strcpy`

String manipulation function with boundary checking that are used incorrectly

- `fgets`

Background: memory regions

- Each running program is allocated a number of distinct memory regions
- *Stack segment*
 - local variables
 - procedure calls
- *Data segment*
 - global initialized variables (data)
 - global uninitialized variables (bss)
 - dynamic variables (heap)
- *Code (text) segment*
 - program instructions (typically, read-only)



Stack

Usually grows towards smaller memory addresses

- Intel, Motorola, SPARC, MIPS

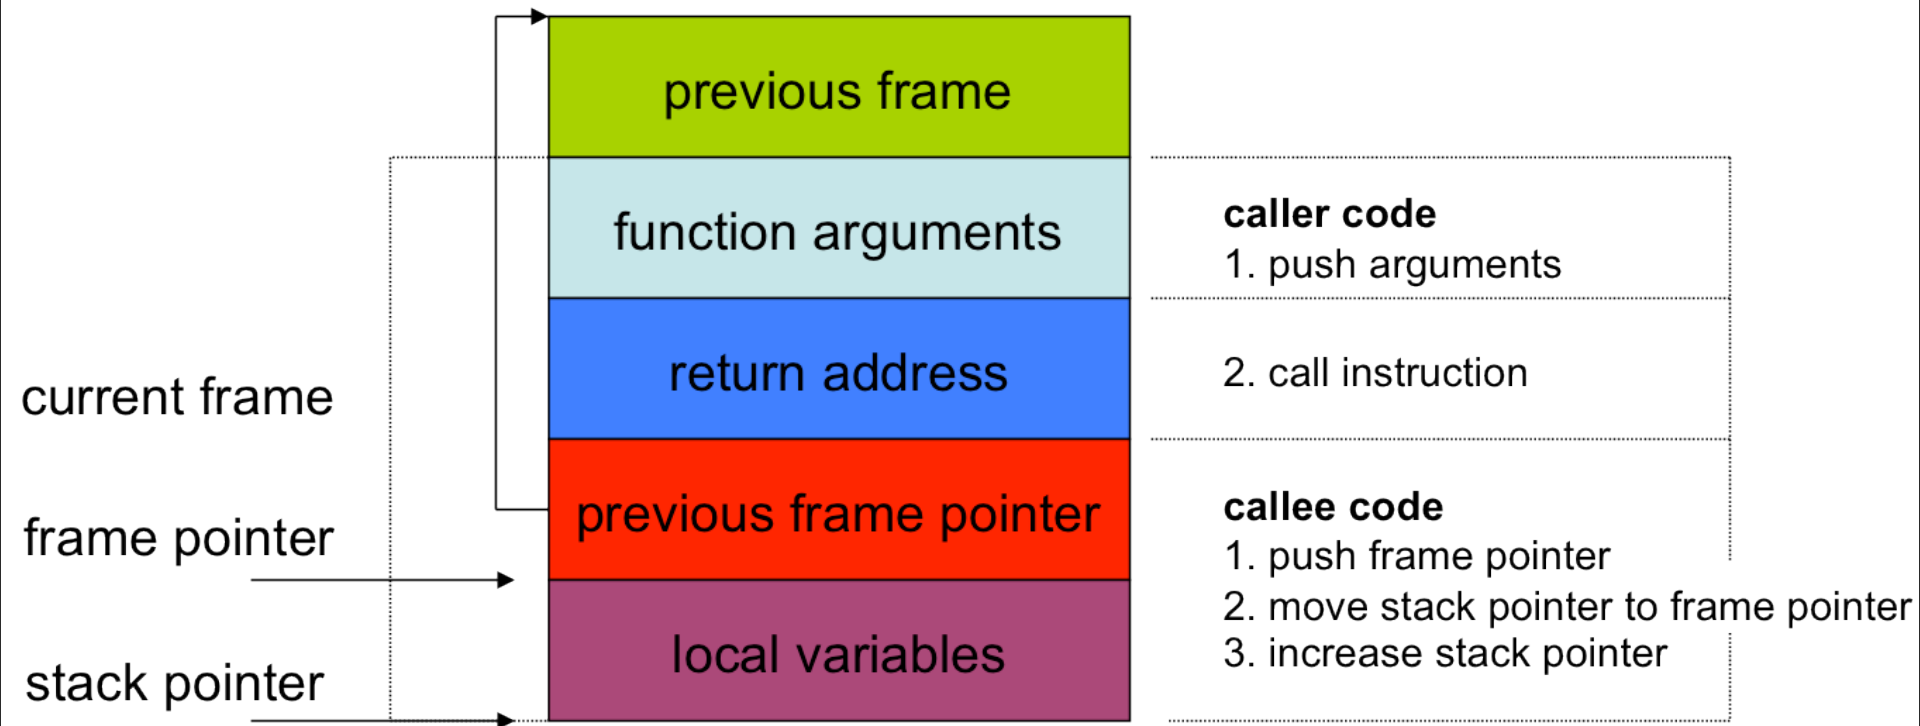
Register points to top of stack

- stack pointer — `sp`
- points to last stack element or first free slot

Composed of *frames*

- pushed on top of stack as consequence of function calls
- address of current frame stored in processor register
- frame/base pointer — `fp`
- used to reference local variables

Stack



Procedure call

```
#include <stdio.h>
```

```
int foo(int a, int b) {  
→   int c = 42;
```

```
    return a * b + c;
```

```
}
```

```
int main(int argc, char **argv) {  
    int c = 0;
```

```
    c = foo(1, 2);
```

```
    printf("%d\n", c);
```

```
    return 0;
```

```
}
```

2
1
Saved IP
Saved BP
42

Under the hood

(gdb) disass main

Dump of assembler code for function main:

```
0x080483fd <+0>:      push    %ebp
0x080483fe <+1>:      mov     %esp,%ebp
0x08048400 <+3>:      and     $0xfffffffff0,%esp
0x08048403 <+6>:      sub     $0x20,%esp
0x08048406 <+9>:      movl    $0x0,0x1c(%esp)
0x0804840e <+17>:     movl    $0x2,0x4(%esp)
0x08048416 <+25>:     movl    $0x1,(%esp)
0x0804841d <+32>:     call    0x80483e4 <foo>
0x08048422 <+37>:     mov     %eax,0x1c(%esp)
0x08048426 <+41>:     mov     $0x8048520,%eax
0x0804842b <+46>:     mov     0x1c(%esp),%edx
0x0804842f <+50>:     mov     %edx,0x4(%esp)
0x08048433 <+54>:     mov     %eax,(%esp)
0x08048436 <+57>:     call    0x8048300
<printf@plt>
0x0804843b <+62>:     mov     $0x0,%eax
0x08048440 <+67>:     leave
0x08048441 <+68>:     ret
End of assembler dump.
```

0x2
0x1
0x8048422

Under the hood

(gdb) disass foo

Dump of assembler code for function foo:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp,%ebp
0x080483e7 <+3>:      sub     $0x10,%esp
0x080483ea <+6>:      movl    $0x2a,-0x4(%ebp)
0x080483f1 <+13>:     mov     0x8(%ebp),%eax
0x080483f4 <+16>:     imul    0xc(%ebp),%eax
0x080483f8 <+20>:     add     -0x4(%ebp),%eax
0x080483fb <+23>:     leave
0x080483fc <+24>:     ret
```

End of assembler dump.

0x2
0x1
0x8048422
0xbffff1a8
0x2a

Let's run it

```
(gdb) br foo
Breakpoint 1 at 0x80483ea: file procedure.c, line 4.
(gdb) r
Starting program: procedure foo

Breakpoint 1, foo (a=1, b=2) at procedure.c:4
4          int c = 42;
(gdb) stepi
6          return a * b + c;
(gdb) x/8wx $ebp -4
0xbffff174: 0x0000002a 0xbffff1a8 0x08048422 0x00000001
0xbffff184: 0x00000002 0x002d8ff4 0x00166225 0x0011f270
```

0x2
0x1
0x8048422
0xbffff1a8
0x2a

Returning

```
(gdb) x/i $eip
```

```
=> 0x80483fb <foo+23>:  leave
```

```
(gdb) info r ebp esp
```

```
ebp                0xbffff178      0xbffff178
```

```
esp                0xbffff168      0xbffff168
```

```
(gdb) si
```

```
(gdb) info r ebp esp
```

```
ebp                0xbffff1a8      0xbffff1a8
```

```
esp                0xbffff17c      0xbffff17c
```

```
(gdb) x/i $eip
```

```
=> 0x80483fc <foo+24>:  ret
```

```
(gdb) si
```

```
0x08048422 in main (argc=2, argv=0xbffff244) at  
procedure.c:12
```

```
(gdb) info r ebp esp eip
```

```
ebp                0xbffff1a8      0xbffff1a8
```

```
esp                0xbffff180      0xbffff180
```

```
eip                0x8048422      0x8048422 <main+37>
```

0x2
0x1

Reference

- [GDB: The GNU Project Debugger](#)
- [Intel 64 and IA-32 Architectures Software Developer's Manual](#)

Vulnerable program

```
#include <stdio.h>
#include <string.h>

void vulnerable(char* param)
{
    char buffer[100];
    strcpy(buffer, param);
}

int main(int argc, char** argv)
{
    vulnerable(argv[1]);
    printf("OK\n");

    return 0;
}
```

Let's crash

```
06-buffer-overflows$ ./vuln test
```

```
OK
```

```
marco@ubuntu:$ ./vuln
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAA
```

```
Segmentation fault (core dumped)
```

Let's crash

```
$ gdb vuln
(gdb) r
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAA
Starting program: vuln
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAA
Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
(gdb) info r esp ebp eip
esp                0xbffff0f0      0xbffff0f0
ebp                0x41414141      0x41414141
eip                0x41414141      0x41414141
(gdb) x/20wx $esp -32
0xbffff0d0:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff0e0:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff0f0:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff100:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff110:      0x41414141      0x41414141      0x41414141      0x41414141
```


To reproduce

Disable defense mechanisms!

```
$ gcc -fno-stack-protector \  
    -U_FORTIFY_SOURCE -D_FORTIFY_SOURCE=0 \  
    vuln.c -o vuln
```

(We'll see these in detail in a few lectures.)

Smashing the stack

Key idea: Overwrite a pointer with the address of code we want to execute

1. locate a pointer that (eventually) will be copied to the EIP register or that points to data that will be copied to the EIP
 - function pointers (on the stack, heap, BSS, ...)
 - saved EBP
 - procedure return address
 - entry in the GOT
 - jmp_buf
2. overwrite the pointer with our value

Smashing the stack

- A procedure contains a local buffer variable allocated on the stack
- The procedure copies user-controlled data to the buffer without verifying that the data size is smaller than the buffer
- The user data overwrites the other variables on the stack, up to the *return address* saved in the function frame
- When the procedure returns, the program fetches the return address from the stack and copies it to the EIP register
- Since we can control the return address, we can jump to an address of our choice

Where do we jump to?

Address inside a buffer whose content is under user control

- PRO: works for remote attacks
- CON: the attacker needs to know the address of the buffer
- CON: the memory page containing the buffer must be executable

Where do we jump to?

Address of an environment variable

- PRO: easy to implement, works with tiny buffers
- CON: only for local exploits
- CON: some programs clean the environment
- CON: the environment must be executable

Where do we jump to?

- Address of a function inside the program
- PRO: works for remote attacks
- PRO: does not require an executable stack
- CON: need to find the right code
- CON: one or more fake frames must be put on the stack


Jumping into the buffer

- The buffer that we are overflowing is usually a good place to put the code (shellcode) that we want to execute
- The buffer is somewhere on the stack, but in most cases the exact address is unknown
- The address must be precise: jumping one byte before or after would just make the application crash
- On the local system, it is possible to calculate the address with a debugger, but it is very unlikely to be the same address on a different machine: any change to the environment variables affects the stack position

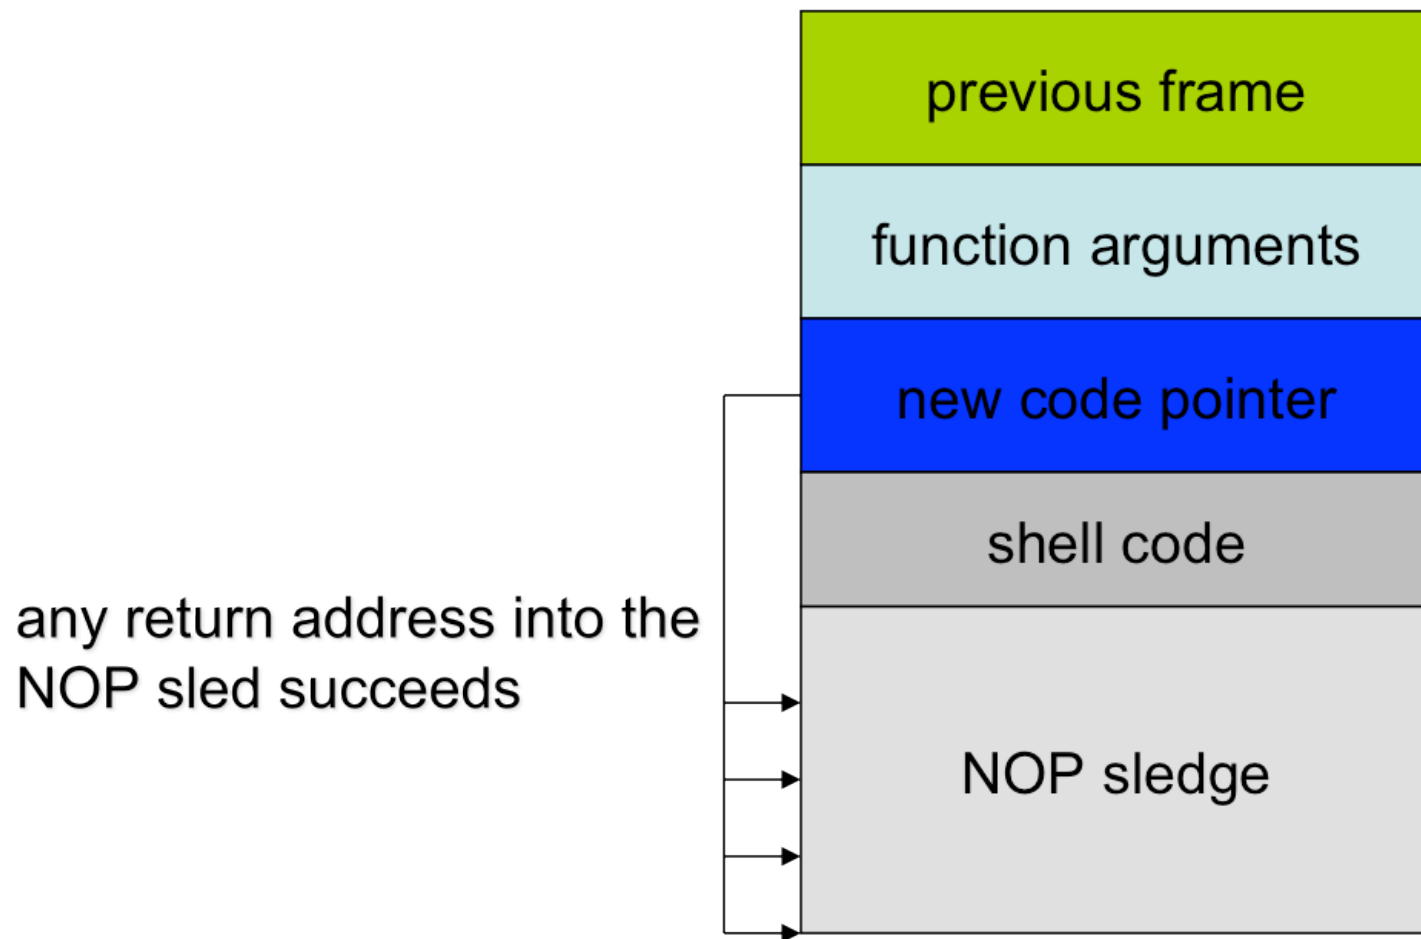
Solution #1: NOP sled

- A sled is a “landing area” in front of the shellcode
 - work arounds the problem of finding the right shellcode address
- Must be created in a way such that whenever the program jumps into it:
 - it always finds a valid instruction
 - it always reaches the end of the sled and the beginning of the shellcode
- The simplest sled is a sequence of NOP instructions
 - single byte instruction (0x90) that does not do anything
 - more complex sleds possible

Assembling the malicious input



Assembling the malicious input



Solution #2: register-based jump

- Find a register that points to the buffer (or somewhere into it)
 - ESP
 - EAX (return value of a function call)
- Locate an instruction that jumps/calls using that register
 - can also be in one of the libraries
 - does not need to be a real instruction, just look for the right sequence of bytes (`jmp *$esp = 0xFF 0xE4`)
- Overwrite the return address with the address of that instruction

(We will talk about address space randomization)

Solution #3: heap spraying

- We will talk about this later

Some history

[Morris worm](#) (1988): overflow in fingerd

- 6,000 machines infected (10% of the Internet)
- Internet had to be switched off
- CERT is created

[CodeRed](#) (2001): overflow in MS IIS server

- 300,000 machines infected in 14 hours

[SQL Slammer](#) (2003): overflow in MS SQL server

- attack: 1 UDP packet
- 75,000+ machines infected in 10 minutes

In 2003, around 75% of the vulnerabilities were buffer overflows

References

- Aleph One, [Smashing The Stack For Fun And Profit](#), Phrack 49, 1996
- spoonm, [Recent Shellcode Developments](#), ReCon 2005

Next time

Introduction to assembly

Read J. Mason et al., [English Shellcode](#), CCS
2009