

Buffer overflows - cont'd

Secure Programming
Lecture 7

In the news

New Feed Google feedburner

[Edit Feed Details...](#) | [Delete Feed...](#) | [Transfer Feed...](#)

Analyze | **Optimize** | **Publicize** | **Troubleshootize** My Feeds

VIEW

[Tips and Tools](#)

REPORTS

[FeedMedic](#)

PodMedic

PodMedic

PodMedic examines links in each of your feed items. If it finds trouble creating an enclosure, it reports the cause of the trouble. PodMedic examines every link in each posting. Of course, you probably link to things that aren't meant to become enclosures, so you can safely ignore reports on links that point to other web pages. Your current report is always displayed below.

Content Item: "2"

♦ **Link:** [REDACTED]

Content Type reported: text/; charset=UTF-8

Media enclosure not created because: The content type returned from your host server was a text type, not a binary media type. Podcast content should always return a binary media type, such as audio/mpeg for audio, or video/mpeg for video.

♦ **Link:** [REDACTED]

Модальное окно: Подтвердите действие на feedburner.google.com

feedburner.google.com

OK

Submit your content to Google • [Privacy Policy](#) ©2004–2013 Google

<http://paul-axe.blogspot.de/2014/01/two-stories-about-xss-on-google.html>

In the news

includes/media/DjVu.php

```
.. .. @@ -180,9 +180,12 @@
180 180     $srcPath = $image->getLocalRefPath();
181 181     # Use a subshell (brackets) to aggregate stderr from both pipeline commands
182 182     # before redirecting it to the overall stdout. This works in both Linux and Windows XP.
183     - $cmd = '(' . wfEscapeShellArg( $wgDjvuRenderer ) . " -format=ppm -page={$page}" .
184     -     " -size={$params['physicalWidth']}x{$params['physicalHeight']}" .
185     -     wfEscapeShellArg( $srcPath );
183 +     $cmd = '(' . wfEscapeShellArg(
184 +         $wgDjvuRenderer,
185 +         "-format=ppm",
186 +         "-page={$page}",
187 +         "-size={$params['physicalWidth']}x{$params['physicalHeight']}",
188 +         $srcPath );
186 189     if ( $wgDjvuPostProcessor ) {
187 190         $cmd .= " | {$wgDjvuPostProcessor}";
188 191     }
```

includes/media/ImageHandler.php

```
.. .. @@ -93,6 +93,7 @@
93 93     if ( !isset( $params['page'] ) ) {
94 94         $params['page'] = 1;
95 95     } else {
96 +     $params['page'] = intval( $params['page'] );
96 97     if ( $params['page'] > $image->pageCount() ) {
97 98         $params['page'] = $image->pageCount();
98 99     }
```

<https://git.wikimedia.org/commitdiff/mediawiki%2Fcore.git/9e0b52b14537a4238239a6e8a2319fe3838fa409>

Procedure call

```
#include <stdio.h>
```

```
int foo(int a, int b) {  
→   int c = 42;
```

```
    return a * b + c;
```

```
}
```

```
int main(int argc, char **argv) {  
    int c = 0;
```

```
    c = foo(1, 2);
```

```
    printf("%d\n", c);
```

```
    return 0;
```

```
}
```

2
1
Saved IP
Saved BP
42

Under the hood

(gdb) disass main

Dump of assembler code for function main:

```
0x080483fd <+0>:      push    %ebp
0x080483fe <+1>:      mov     %esp,%ebp
0x08048400 <+3>:      and     $0xfffffffff0,%esp
0x08048403 <+6>:      sub     $0x20,%esp
0x08048406 <+9>:      movl    $0x0,0x1c(%esp)
0x0804840e <+17>:     movl    $0x2,0x4(%esp)
0x08048416 <+25>:     movl    $0x1, (%esp)
0x0804841d <+32>:     call    0x80483e4 <foo>
0x08048422 <+37>:     mov     %eax,0x1c(%esp)
0x08048426 <+41>:     mov     $0x8048520,%eax
0x0804842b <+46>:     mov     0x1c(%esp),%edx
0x0804842f <+50>:     mov     %edx,0x4(%esp)
0x08048433 <+54>:     mov     %eax, (%esp)
0x08048436 <+57>:     call    0x8048300
<printf@plt>
0x0804843b <+62>:     mov     $0x0,%eax
0x08048440 <+67>:     leave
0x08048441 <+68>:     ret
End of assembler dump.
```

0x2
0x1
0x8048422

Under the hood

(gdb) disass foo

Dump of assembler code for function foo:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp, %ebp
0x080483e7 <+3>:      sub     $0x10, %esp
0x080483ea <+6>:      movl    $0x2a, -0x4(%ebp)
0x080483f1 <+13>:     mov     0x8(%ebp), %eax
0x080483f4 <+16>:     imul    0xc(%ebp), %eax
0x080483f8 <+20>:     add     -0x4(%ebp), %eax
0x080483fb <+23>:     leave
0x080483fc <+24>:     ret
```

End of assembler dump.

0x2
0x1
0x8048422
0xbffff1a8
0x2a

Let's run it

```
(gdb) br foo
Breakpoint 1 at 0x80483ea: file procedure.c, line 4.
(gdb) r
Starting program: procedure foo

Breakpoint 1, foo (a=1, b=2) at procedure.c:4
4          int c = 42;
(gdb) stepi
6          return a * b + c;
(gdb) x/8wx $ebp -4
0xbffff174: 0x0000002a 0xbffff1a8 0x08048422 0x00000001
0xbffff184: 0x00000002 0x002d8ff4 0x00166225 0x0011f270
```

0x2
0x1
0x8048422
0xbffff1a8
0x2a

Returning

```
(gdb) x/i $eip
```

```
=> 0x80483fb <foo+23>:  leave
```

```
(gdb) info r ebp esp
```

```
ebp                0xbffff178      0xbffff178
```

```
esp                0xbffff168      0xbffff168
```

```
(gdb) si
```

```
(gdb) info r ebp esp
```

```
ebp                0xbffff1a8      0xbffff1a8
```

```
esp                0xbffff17c      0xbffff17c
```

```
(gdb) x/i $eip
```

```
=> 0x80483fc <foo+24>:  ret
```

```
(gdb) si
```

```
0x08048422 in main (argc=2, argv=0xbffff244) at  
procedure.c:12
```

```
(gdb) info r ebp esp eip
```

```
ebp                0xbffff1a8      0xbffff1a8
```

```
esp                0xbffff180      0xbffff180
```

```
eip                0x8048422      0x8048422 <main+37>
```

0x2
0x1

BUFFER OVERFLOWS

Vulnerable program

```
#include <stdio.h>
#include <string.h>

void vulnerable(char* param)
{
    char buffer[100];
    strcpy(buffer, param);
}

int main(int argc, char** argv)
{
    vulnerable(argv[1]);
    printf("OK\n");

    return 0;
}
```

Let's crash

```
06-buffer-overflows$ ./vuln test
```

```
OK
```

```
marco@ubuntu:$ ./vuln
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAA
```

```
Segmentation fault (core dumped)
```

Let's crash

```
$ gdb vuln
(gdb) r
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAA
Starting program: vuln
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAA
Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
(gdb) info r esp ebp eip
esp                0xbffff0f0      0xbffff0f0
ebp                0x41414141      0x41414141
eip                0x41414141      0x41414141
(gdb) x/20wx $esp -32
0xbffff0d0:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff0e0:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff0f0:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff100:      0x41414141      0x41414141      0x41414141      0x41414141
0xbffff110:      0x41414141      0x41414141      0x41414141      0x41414141
```

To reproduce

Disable defense mechanisms!

```
$ gcc -fno-stack-protector \  
    -U_FORTIFY_SOURCE -D_FORTIFY_SOURCE=0 \  
    vuln.c -o vuln
```

(We'll see these in detail in a few lectures.)

Smashing the stack

Key idea: Overwrite a pointer with the address of code we want to execute

1. locate a pointer that (eventually) will be copied to the EIP register or that points to data that will be copied to the EIP
 - function pointers (on the stack, heap, BSS, ...)
 - saved EBP
 - procedure return address
 - entry in the GOT
 - jmp_buf
2. overwrite the pointer with our value

Smashing the stack

- A procedure contains a local buffer variable allocated on the stack
- The procedure copies user-controlled data to the buffer without verifying that the data size is smaller than the buffer
- The user data overwrites the other variables on the stack, up to the *return address* saved in the function frame
- When the procedure returns, the program fetches the return address from the stack and copies it to the EIP register
- Since we can control the return address, we can jump to an address of our choice

Where do we jump to?

Address inside a buffer whose content is under user control

- PRO: works for remote attacks
- CON: the attacker needs to know the address of the buffer
- CON: the memory page containing the buffer must be executable

Where do we jump to?

Address of an environment variable

- PRO: easy to implement, works with tiny buffers
- CON: only for local exploits
- CON: some programs clean the environment
- CON: the environment must be executable

Where do we jump to?

- Address of a function inside the program
- PRO: works for remote attacks
- PRO: does not require an executable stack
- CON: need to find the right code
- CON: one or more fake frames must be put on the stack

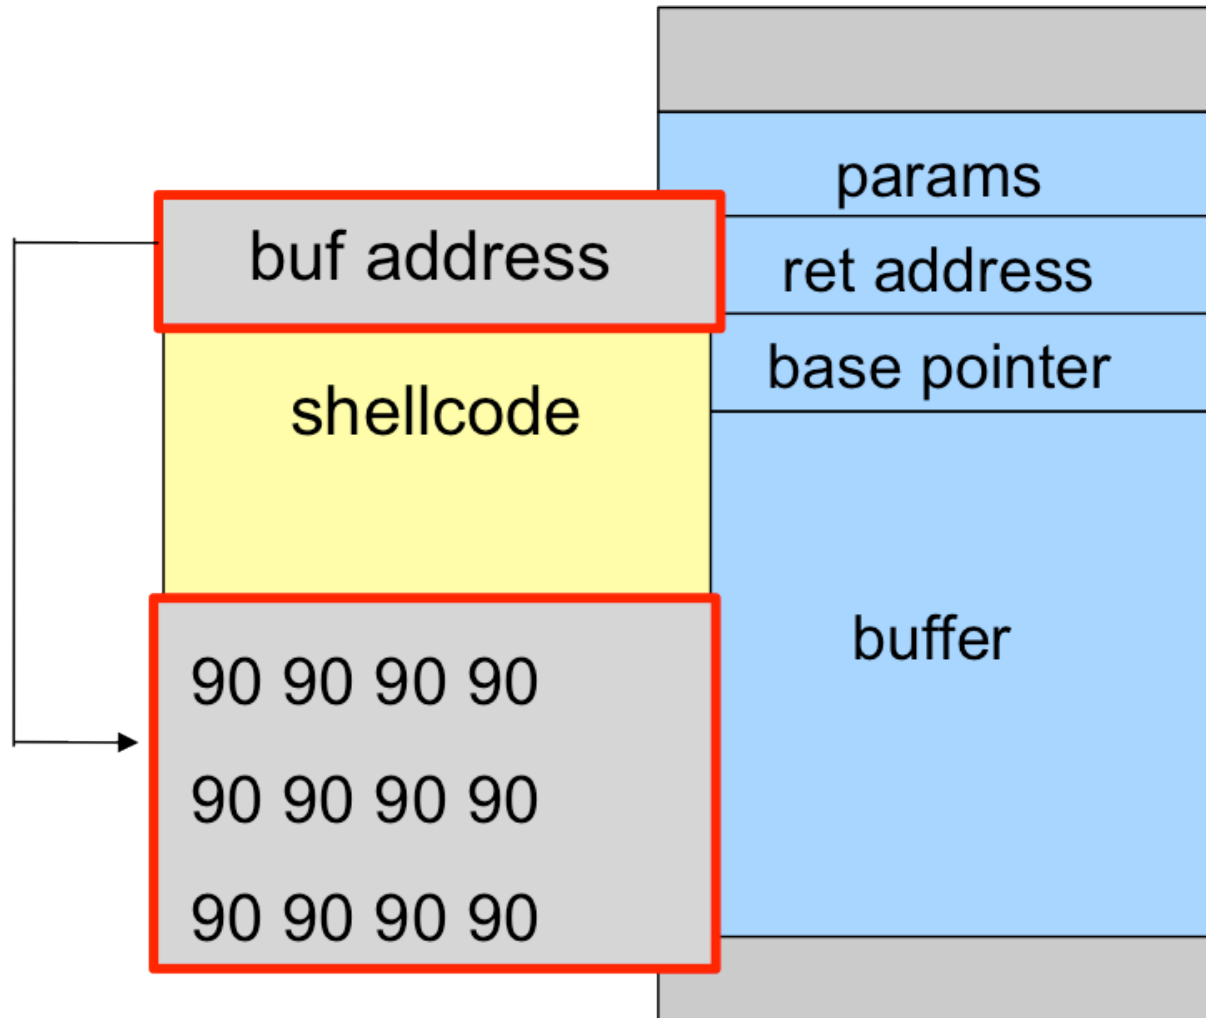
Jumping into the buffer

- The buffer that we are overflowing is usually a good place to put the code (shellcode) that we want to execute
- The buffer is somewhere on the stack, but in most cases the exact address is unknown
- The address must be precise: jumping one byte before or after would just make the application crash
- On the local system, it is possible to calculate the address with a debugger, but it is very unlikely to be the same address on a different machine: any change to the environment variables affects the stack position


Solution #1: NOP sled

- A sled is a “landing area” in front of the shellcode
 - work arounds the problem of finding the right shellcode address
- Must be created in a way such that whenever the program jumps into it:
 - it always finds a valid instruction
 - it always reaches the end of the sled and the beginning of the shellcode
- The simplest sled is a sequence of NOP instructions
 - single byte instruction (0x90) that does not do anything
 - more complex sleds possible

Assembling the malicious input



Assembling the malicious input



Solution #2: register-based jump

- Find a register that points to the buffer (or somewhere into it)
 - ESP
 - EAX (return value of a function call)
- Locate an instruction that jumps/calls using that register
 - can also be in one of the libraries
 - does not need to be a real instruction, just look for the right sequence of bytes (`jmp *$esp = 0xFF 0xE4`)
- Overwrite the return address with the address of that instruction

(We will talk about address space randomization)

Solution #3: heap spraying

- We will talk about this later

Some history

[Morris worm](#) (1988): overflow in fingerd

- 6,000 machines infected (10% of the Internet)
- Internet had to be switched off
- CERT is created

[CodeRed](#) (2001): overflow in MS IIS server

- 300,000 machines infected in 14 hours

[SQL Slammer](#) (2003): overflow in MS SQL server

- attack: 1 UDP packet
- 75,000+ machines infected in 10 minutes

In 2003, around 75% of the vulnerabilities were buffer overflows

x86

ASSEMBLY REFRESHER

Registers


General purpose registers:

- AX (accumulator)
- BX (base)
- CX (counter)
- DX (data)
- SP (stack pointer)
- BP (stack base pointer)
- SI (source index)
- DI (destination index)

- Can be accessed in 5 modes: 64, 32, 16, 8 bit (LSB and MSB)

RAX			
	EAX		
		AX	
		AH	AL

Flags register



- S Indicates a Status Flag
- C Indicates a Control Flag
- X Indicates a System Flag

☐ Reserved bit positions. DO NOT USE.
 Always set to values previously read.

Instruction pointer register

The instruction pointer (EIP) register


- Contains the offset in the current code segment for the next instruction to be executed

Cannot be accessed directly by software:

- It is advanced from one instruction boundary to the next in straight-line code
- moved ahead or backwards by a number of instructions when executing JMP, Jcc, CALL, RET, and IRET instructions.

Endiannes

- x86 is a little-endian architecture: the least significant byte goes into the lowest value slot



Data transfer instructions

- `mov src, dest`

Copy the source operand (immediate, register, memory) into the destination operand (register, memory)

- `lea src, dest`

Calculate the address of the source operand (immediate, register, memory) and load it into the destination operand (register)

(Used also for general-purpose arithmetic operations)

Control flow instructions

- `test arg1, arg2`
Perform a bit-wise logical AND on `arg1`, `arg2` and set ZF (zero), SF (sign), and PF (parity) flags
- `cmp arg2, arg1`
Perform a signed subtraction of `arg2` from `arg1` and set ZF, SF, PF accordingly
- `jmp loc`
Unconditional jump
- `jcc loc`
Conditional jump (`je`, `jne`, `jb`, `jbe`, `ja`, `jae`, `jl`, `jle`, `jz`, `jnz`, etc.)

Function calls

- `call proc`

Push the address of the next opcode on the stack and jump to specified address

- `ret [val]`

Load the next value on the stack into EIP, pop the specified number of bytes off the stack (none if not specified)

Arithmetic instructions

- `add src, dst`
Add the source operand to destination operand and store the result in the destination operand
- `sub src, dst`
- `mul arg`
Multiply `arg` by the value of the corresponding byte-length in `AX`
- `div arg`
- `inc arg`
- `dec arg`
- ...

Logic instructions

- `and src, dest`

Perform a bit-wise AND of the operands and store the result in the destination operand.

- `or src, dest`
- `xor src, dest`
- `not arg`

Perform a bit-wise inversion of arg

Stack instructions

- `push arg`

Decrement the stack pointer and store the value of the argument in the location pointed to by SP

- `pop arg`

Load the data stored in the location pointed to by SP into the argument and then increments the SP

References

- [GDB: The GNU Project Debugger](#)
- [Intel 64 and IA-32 Architectures Software Developer's Manual](#)
- Aleph One, [Smashing The Stack For Fun And Profit](#), Phrack 49, 1996
- spoonm, [Recent Shellcode Developments](#), ReCon 2005