

Heap Overflows

Secure Programming
Lecture 13

Announcement

- Homework 2 has been graded
 - Let me know of any problems, issues, things that are unclear
- Homework 3 deadline is coming up
Tuesday March 11, 11:59pm
 - Don't procrastinate



JORGE CHAM © 2010

WWW.PHDCOMICS.COM

title: "Procrascorelation" - originally published 10/27/2010

In the news

```
static OSStatus
SSLVerifySignedServerKeyExchange(SSLContext *ctx, bool isRsa, SSLBuffer signedParams,
                                uint8_t *signature, UInt16 signatureLen)
{
    OSStatus      err;
    ...

    if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
        goto fail;
    ...

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}
```

<https://www.imperialviolet.org/2014/02/22/applebug.html>

Heap basics

- Region for dynamically allocated memory
 - Dynamic allocation
 - Dynamic de-allocation (freeing)
- Operating system
 - `sbrk()`: allows to change the “program break”, i.e., the end of the process’ data segment
 - `mmap()`
- Heap manager (“allocator”)
 - Handles memory requests from user programs
 - Acts as interface to the `brk()` system call
 - Standard `malloc()`, `free()`, `realloc()`, `calloc()` functions
 - Different implementation in different systems

Heap management implementations

Algorithm	Operating system
<i>Doug Lea's ptmalloc</i>	<i>Linux glibc</i>
System V	Solaris
<i>BSD phk</i>	<i>*BSD</i>
RtlHeap	Windows

Heap basic

```
#define MAX_PATH_LEN 1024

void handle_request(struct req *r) {
    char *s = (char *) malloc(MAX_PATH_LEN);
    strncpy(s, r->path, MAX_PATH_LEN - 1);
    s[MAX_PATH_LEN - 1] = '\0';

    /* do something with s */

    free(s);
}
```

Heap overflows

```
#define MAX_PATH_LEN 1024

void handle_request(struct req *r) {
    char *s = (char *) malloc(MAX_PATH_LEN);
    strcpy(s, r->path);
    s[MAX_PATH_LEN - 1] = '\0';

    /* do something with s */

    free(s);
}
```

Naive malloc

- Let's try to think how we could implement a memory allocator

Common design and requirements

- Common goals
 - Efficiency (time): malloc, free, realloc should be fast
 - Efficiency (space): don't waste space
- Book-keeping information (metadata):
 - Pointers, size values, indexes to regions
 - Commonly done in-band, i.e., along-side the actual data values (for space efficiency)
- Coalescing of unused regions
 - Merge two or more free pieces of memory that are next to each other (avoid fragmentation)

Heap overflow vulnerabilities

- Overflow of data allocated on the heap
 - Allocate some memory on the heap (buffer, structure, etc.)
 - Write into this region exceeding its boundaries
- Successful exploitation of heap overflows was [demonstrated](#) by Solar Designer (2000)
- Standard reference
MaXX, [Vudo malloc tricks](#), 2001
- General idea of such attacks: take advantage of the mixing of data and control information on the heap to overwrite critical control structure

Doug Lea's malloc

- Memory allocator in glibc
 - Description: <http://g.oswego.edu/dl/html/malloc.html>
 - Source: <http://code.metager.de/source/xref/glibc/malloc/>
- Key functionalities
 - Chunk management
 - Bin management
 - Memory allocation
 - Memory deallocation
 - List handling

Heap memory layout



- Heap is divided into contiguous *chunks* of memory
 - Chunks can be allocated, freed, split, combined
- No two free chunks may be physically adjacent
 - Except “fastbins” (bins containing small chunks < 64 bytes)
- Coalescing via boundary tags design (with space optimizations)
- Best-fit via binning

malloc_chunk

```
struct malloc_chunk {  
    size_t prev_size;  
    size_t size;  
    struct malloc_chunk *fd;  
    struct malloc_chunk *bk  
}
```

- “Misleading declaration”
- Use it as a “view” into memory to access fields at known offsets
- Fields have different interpretation depending on this and previous chunk state (allocated or not)

malloc_chunk

```
struct malloc_chunk {  
    size_t prev_size;  
    size_t size;  
    struct malloc_chunk *fd;  
    struct malloc_chunk *bk  
}
```

If the previous chunk is free, it contains its size.

Otherwise, it holds user data of the previous chunk (reduce waste)

malloc_chunk

```
struct malloc_chunk {  
    size_t prev_size;  
    size_t size;  
    struct malloc_chunk *fd;  
    struct malloc_chunk *bk  
}
```

Holds the chunk size in bytes.

size = requested memory
+ 8 bytes (size_t * 2)
- 4 bytes (prev_size
of next

chunk)

Always multiple of 8
(alignment). Last 3 bits of
size are logically 0.

Reuse last bits:

P: PREV_INUSE (0x1)

M: IS_MMAPPED (0x2)

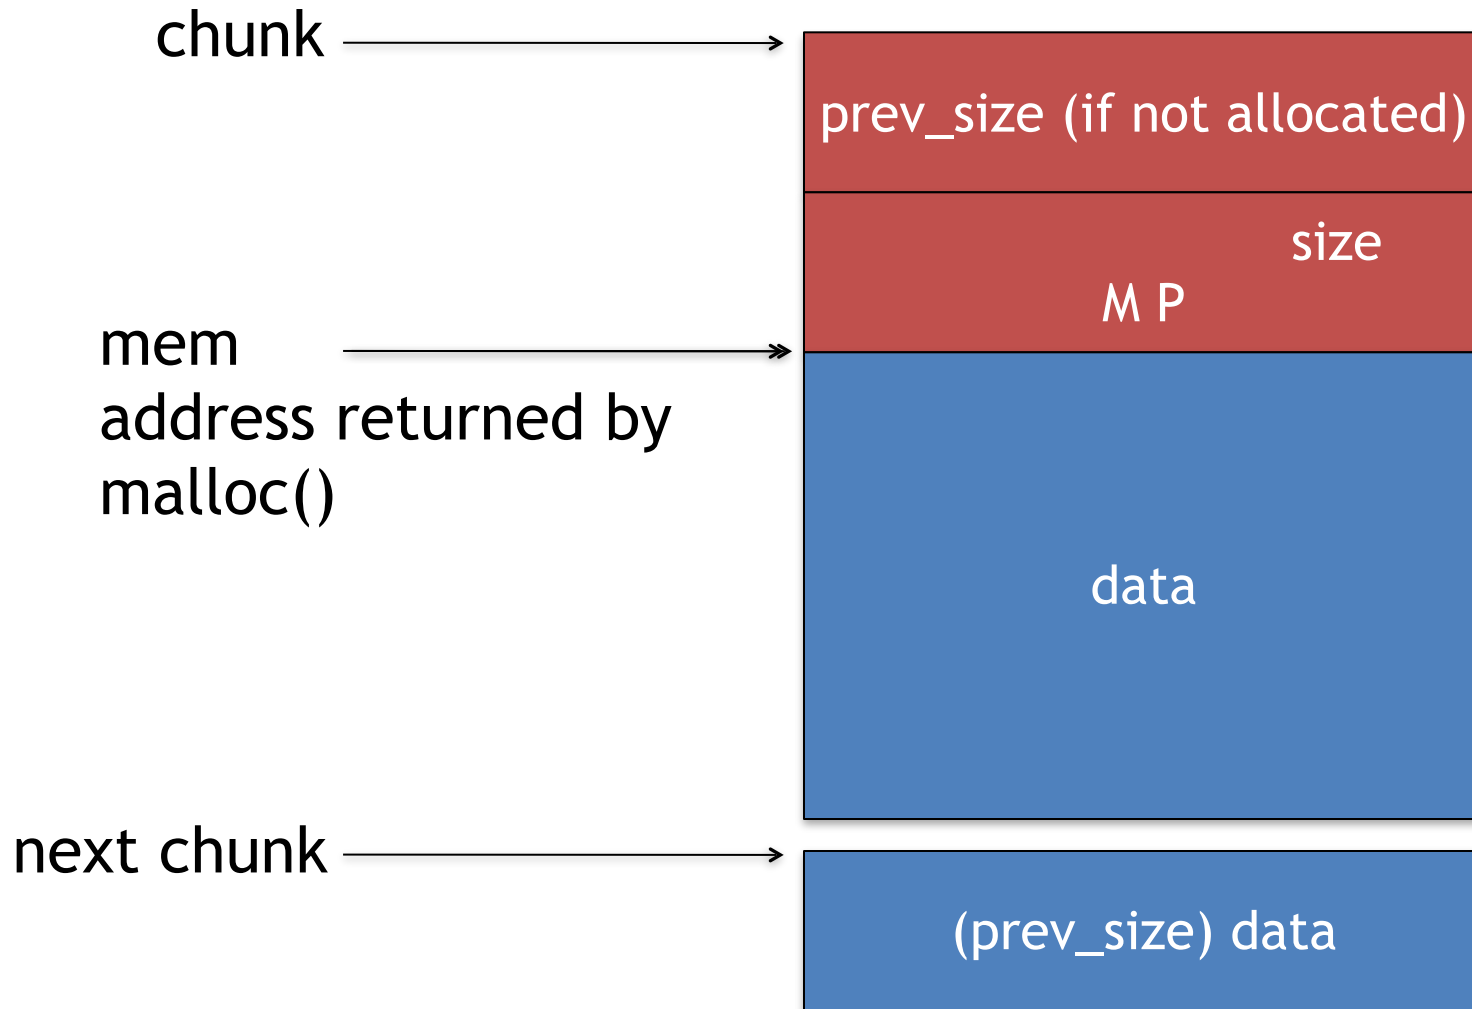
malloc_chunk

```
struct malloc_chunk {  
    size_t prev_size;  
    size_t size;  
    struct malloc_chunk *fd; —  
    struct malloc_chunk *bk  
}
```

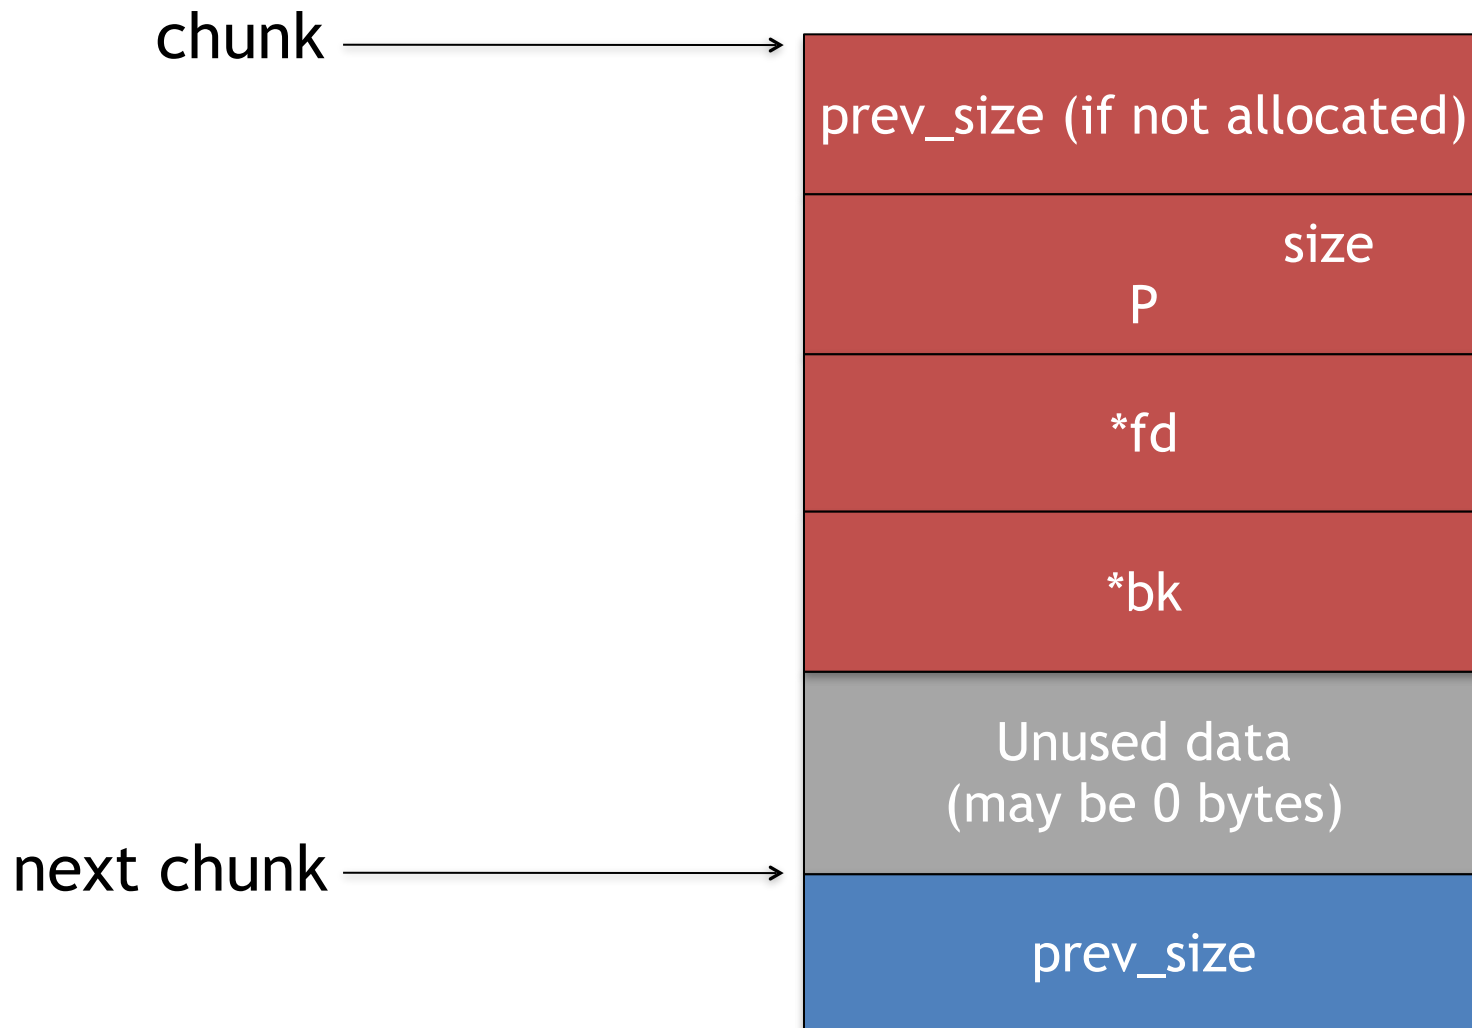
If chunk is free, they contain pointers to next and previous free blocks (doubly-linked list)

If chunk is used, they contain user data.


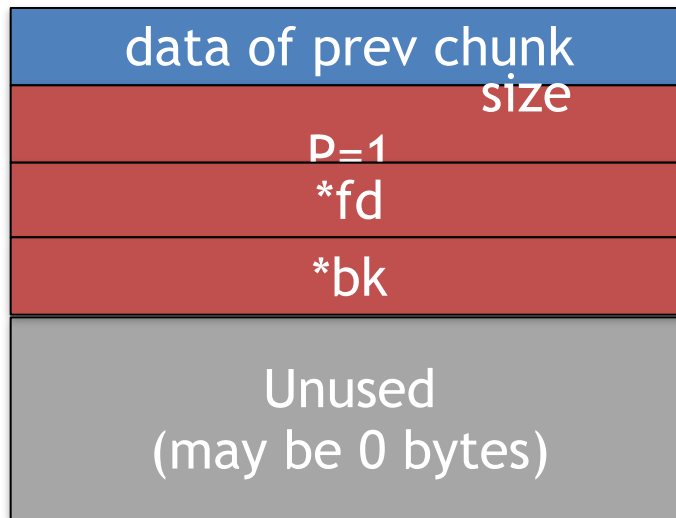

Allocated chunk



Free chunk

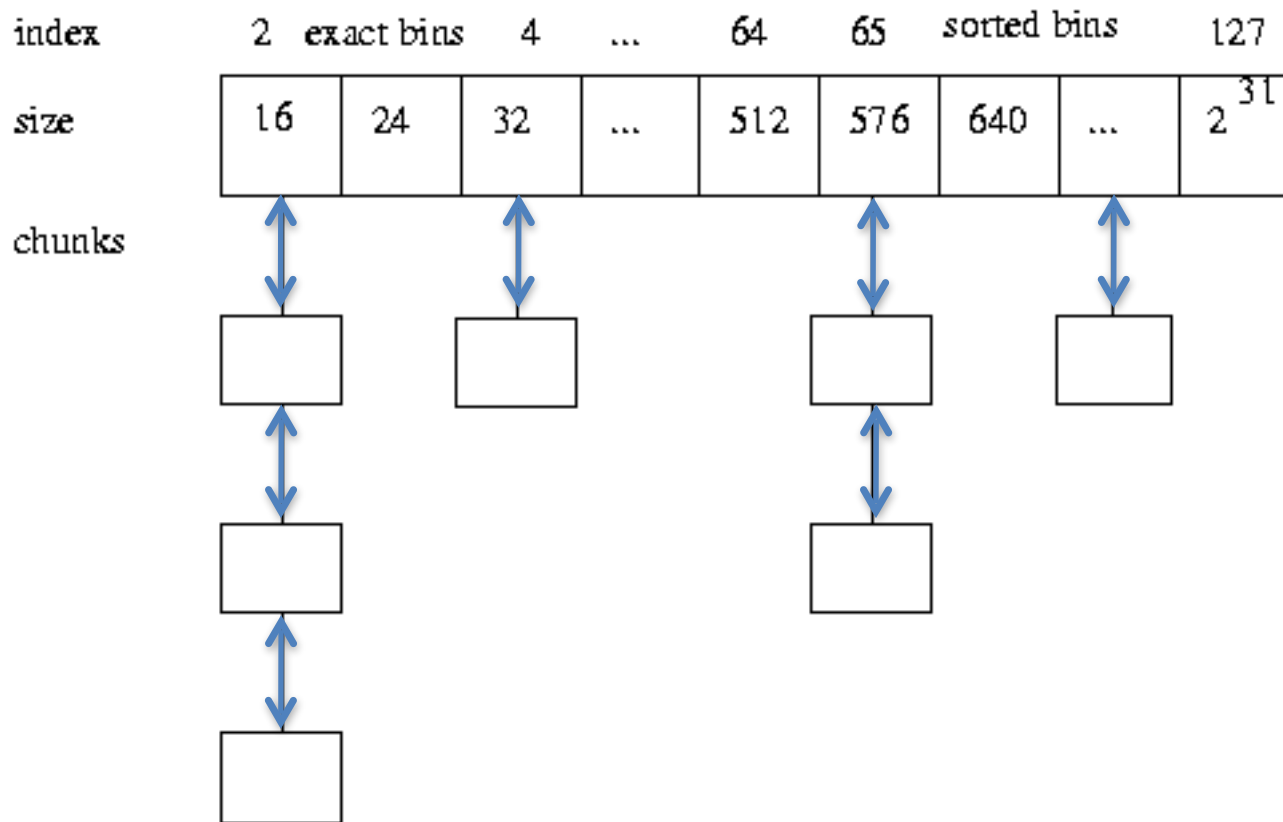


Chunks



Bin management

- Available chunks are maintained in bins grouped by size




List handling: unlink()

- When chunks are allocated or freed, their entries have to be taken off or inserted into the appropriate list
- Macro unlink is used to take off a chunk with its pointers FD and BK

```
#define unlink(P, BK, FD)
{
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```

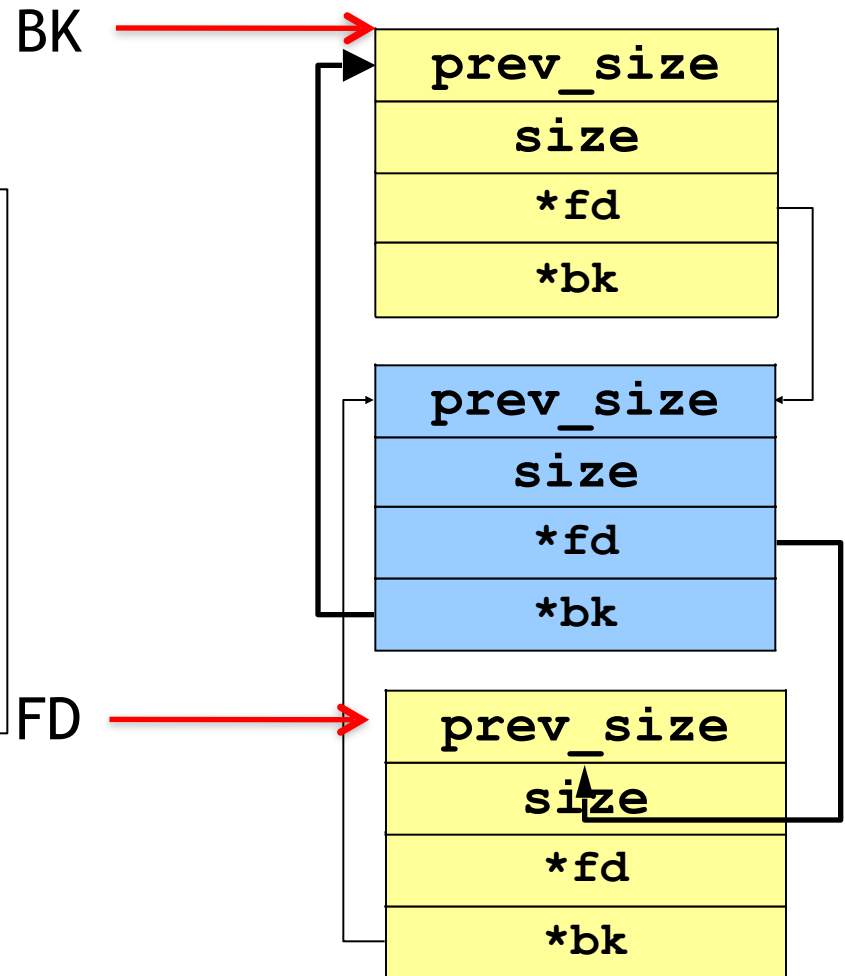
List handling: unlink()

```
#define unlink(P, BK, FD)
{
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```



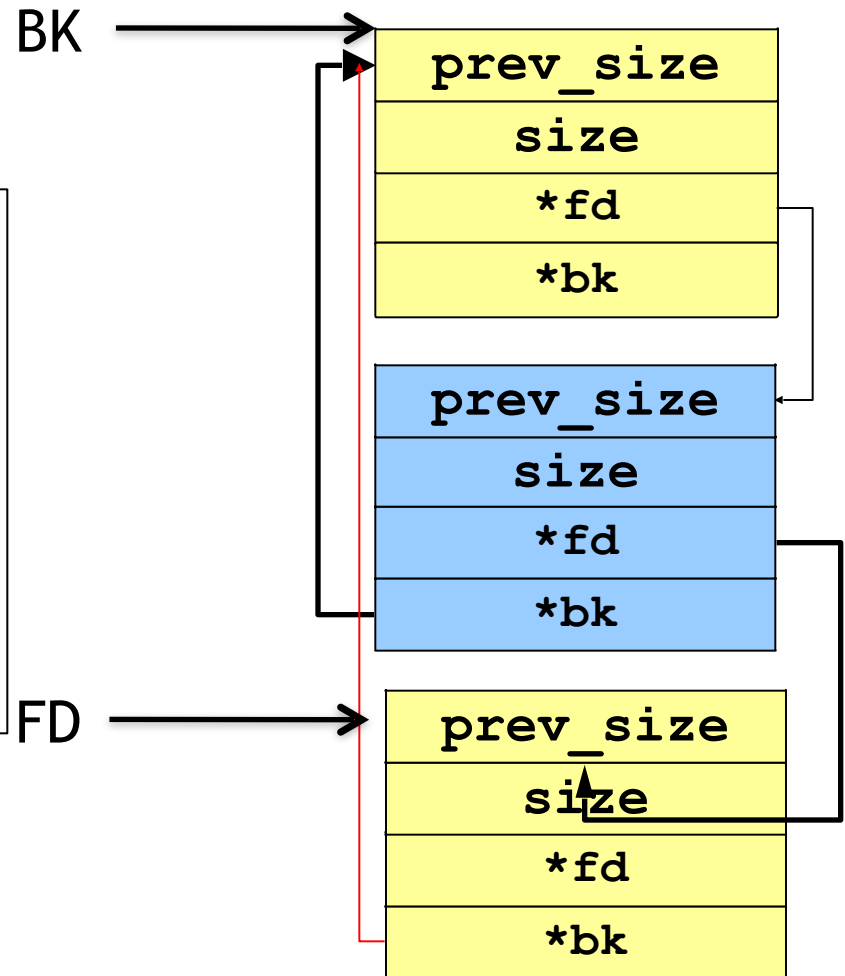
List handling: unlink()

```
#define unlink(P, BK, FD)
{
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```



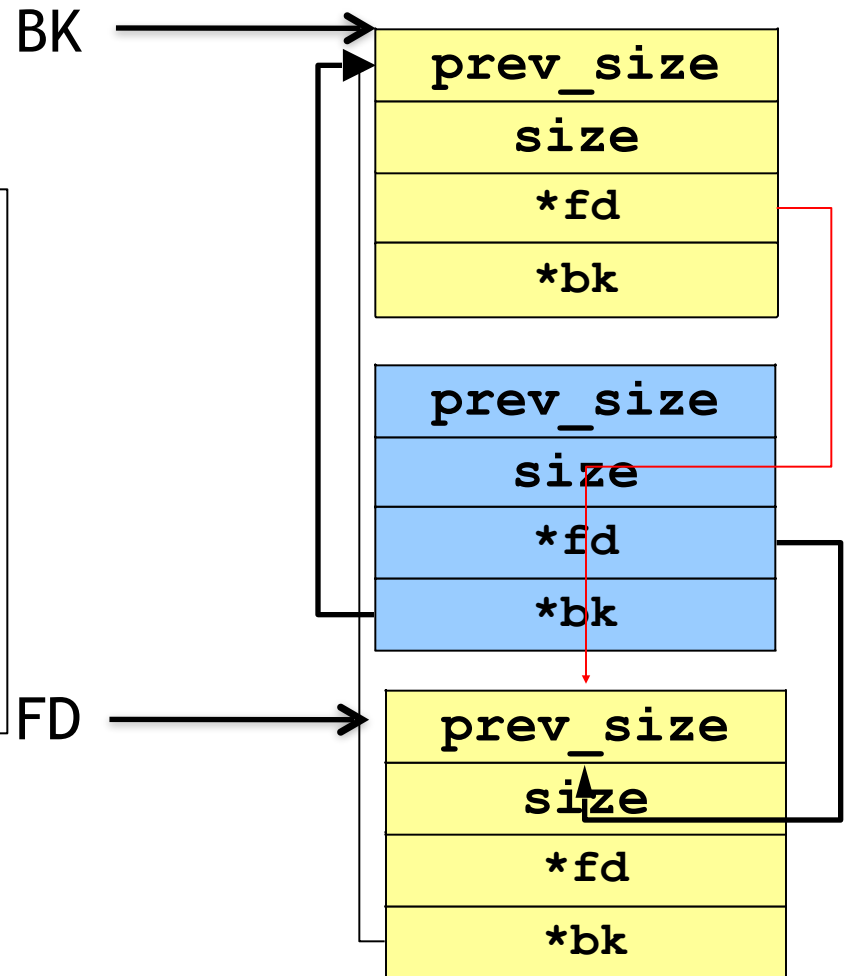
List handling: unlink()

```
#define unlink(P, BK, FD)
{
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```



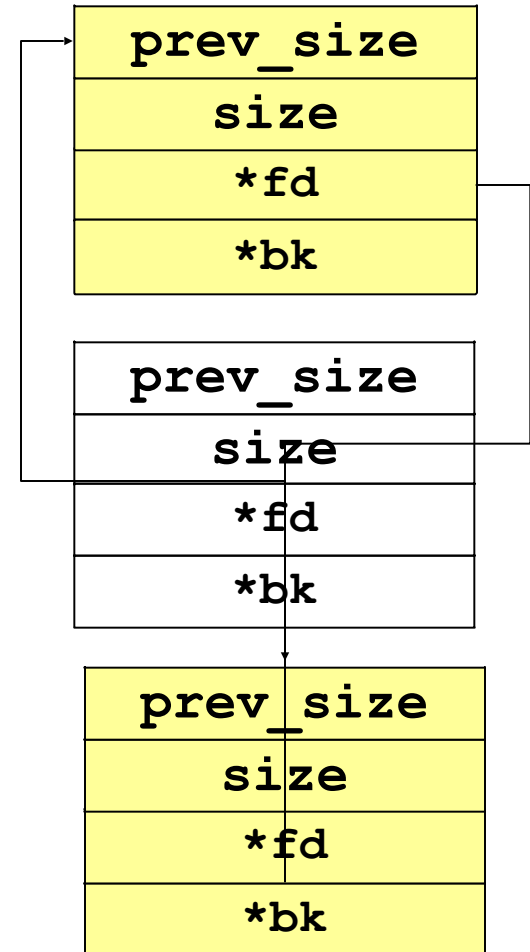
List handling: unlink()

```
#define unlink(P, BK, FD)
{
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```



List handling: unlink()

```
#define unlink(P, BK, FD)
{
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```



Exploiting the `unlink()` macro

- What happens if a chunk is overflown?
 - Control data stored in next chunk is overwritten
 - Attacker can set up a “fake chunk” and abuse `unlink()`

```
P:  struct malloc_chunk {          BK = P->bk;
P+0:  size_t prev_size;          FD = P->fd;
P+4:  size_t size;              FD->bk = BK; // *(P->fd+12) = P->bk
P+8:  struct malloc_chunk *fd;   BK->fd = FD; // *(P->bk+8) = P->fd
P+c:  struct malloc_chunk *bk
      }
```

Overwrite address stored in `FD + 12` with `BK`

Exploiting the unlink() macro


Overwrite an arbitrary memory location with an arbitrary value

- Overwrite a function pointer with address of the shellcode
 - Often, overwrite an entry in the GOT table
 - <http://althing.cs.dartmouth.edu/secref/resources/plt-got.txt>
 - <http://www.technovelty.org/linux/plt-and-got-the-key-to-code-sharing-and-dynamic-libraries.html>
- When function is later invoked, shellcode is executed instead

Exploiting the `unlink()` macro

1. Vulnerable program allocates two adjacent memory chunks, named X and Y
2. The attacker overflows the chunk X creating two fake (free) chunks over Y, called W and Z
3. When X is freed, it will be merged with W and the `unlink()` macro will be called
4. Since W and Z are under the attacker's control, arbitrary values can be specified for their headers

Exploiting the unlink() macro



Exploiting the unlink() macro

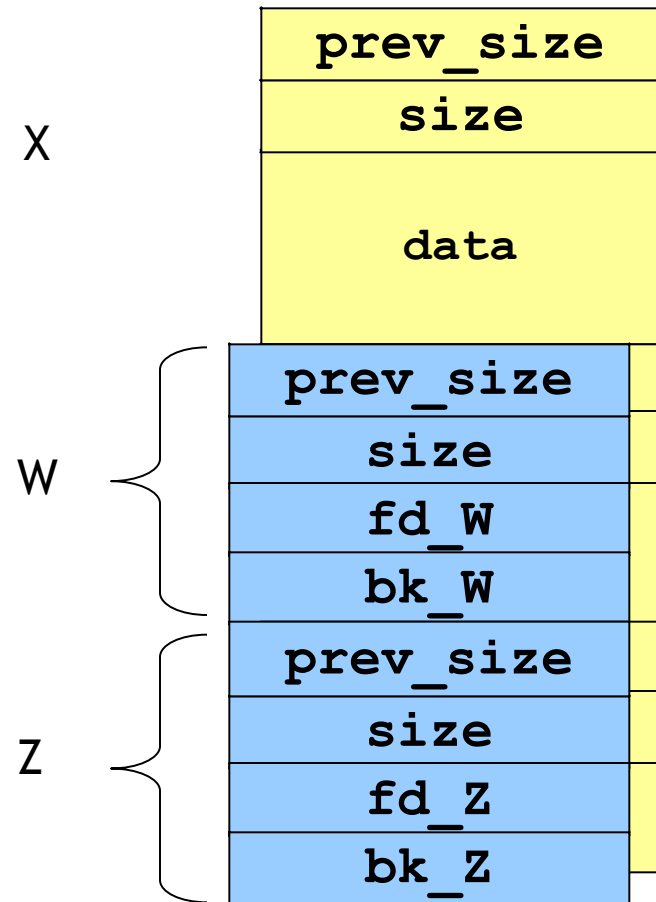
1. free(X) is called
2. W is examined and Z is found using (W + W->size)
 - Z says that W is free
3. unlink(W, fd_w, bk_w) is called

***(fd_w + 12) = bk_w**

***(bk_w + 8) = fd_w**

fd_w must be set to the address to be overwritten - 12

bk_w must be set to the value to be written (shellcode)

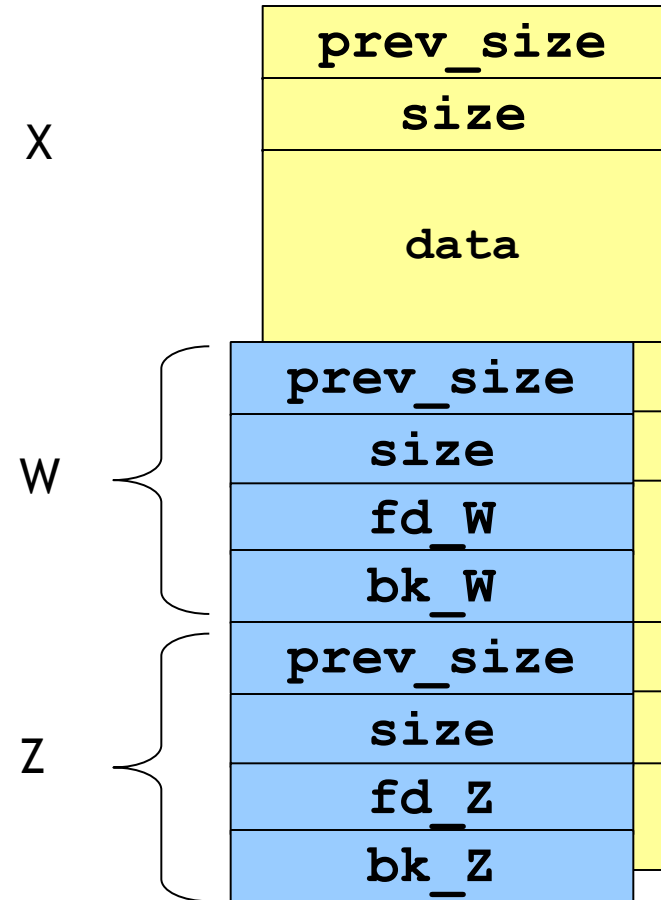


Setting up the buffer

- Find address to overwrite (FUNC_POINTER)
 - E.g., address of GOT entry for free()
- First 8 bytes of X are overwritten by free
 - fd and bk pointers
 - We'll need to skip them
- Unlink modifies the BK + 8
 - We will also need to skip this

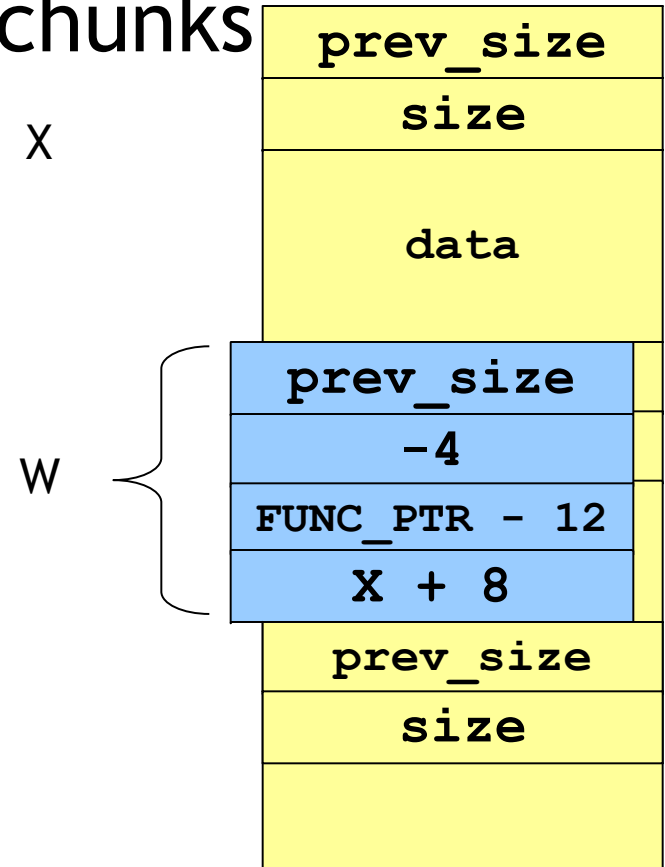
Setting up the buffer

- Target address: 8 bytes into X
 - Skip new fd and bk set by free
- Jump instructions (0xeb 0xc) to jump over the part of the buffer modified by unlink
- Actual shellcode
- Padding to overwrite Y
- Fake chunk (W)
 - fd= FUNC_POINTER - 12
 - bk= X + 8
- Fake chunk (Z)
 - PREV_INUSE = 0



Neat trick

- We don't need two fake chunks
- Fake chunk W
 - size = -4 (0xfffffffffc)



Take away points

- Another example of problems arising from storing control information in-band
 - As when storing return address on the stack
- Exploiting certain vulnerabilities requires a detailed understanding of low-level implementation details