

Integer overflow & Algorithmic complexity vulnerabilities

Lecture 18

Module questionnaire

- Please take a few minutes to complete the module questionnaire

INTEGER VULNERABILITIES

Numbers in a CPU

- Mismatch between machine arithmetic and mathematical arithmetic
- Integer representation in today's architecture
 - Two's complement
 - See section 4.2.1 of [Intel Software Developer's Manual](#)
 - Good introduction:
<http://www.cs.cornell.edu/~tomf/notes/cps104/twoscomp.html>

Two's complement

Unsigned integers

- Ordinary binary values ranging from 0 to max positive number that can be encoded given the integer size (N)
- Range: $[0, 2^N - 1]$
 - 4 bit: [0000, 1111]

Signed integers

- Two's complement representation
- Sign bit
- Range: $[-2^{N-1}, 2^{N-1} - 1]$
 - 4 bit, [1000, 0111]

Integer overflow/underflow

- Integer overflow/underflow
 - Arithmetic expression results in value that is larger/smaller than can be represented on machine type
 - Typically, result silently “wraps around”

```
char * vuln_alloc(int size, int n) {  
    return (char *) malloc(size * n);  
}
```

- If the multiplication overflows, the allocated memory may be smaller than expected, which may lead to a buffer overflow

Width conversion

- Insecure conversion of an integral type to a wider or narrower integral type which has different range of values

```
#include <sys/types.h>
void vuln_copy(int16_t n, char *p, char *q) {
    uint32_t m = n;
    memcpy(p, q, m);
}
```

Signed/Unsigned conversion

- Converting a signed integer type to an unsigned integer type of the *same width* can change a negative number to a large positive number (or vice versa)

```
void vuln_copy(int n, char *p, char *q) {  
    if (n > 800) // too large  
        return;  
    memcpy(p, q, m);  
}
```

n = -1; m = 4294967295

- If n is negative, it will pass the bound check, but will copy a large number of bytes when promoted to unsigned int

Quiz

- `short x = 0x7FFF; x++;`
 - What is the value of x?
- `unsigned long l; short x = -2;`
`l = x;`
 - What is the value of l?

Exploitable integer overflows

```
int main(int argc, char *argv[])
{
    char buf[512];
    long max;
    short len;

    max = sizeof(buf);
    len = strlen(argv[1]);
    printf("max %ld len %hd\n", max, len);
    if (len < max) {
        strcpy(buf, argv[1]);
    }

    return 0;
}
```

Integer Overflows

```
$ ./integeroverflow `python -c 'print "A" * 0x7fff`  
max 512 len 32767
```

```
$ ./integeroverflow `python -c 'print "A" * 0x8000`  
max 512 len -32768  
Segmentation fault (core dumped)
```

Binary search in JDK

```
1: public static int binarySearch(int[] a, int key) {
2:     int low = 0;
3:     int high = a.length - 1;
4:
5:     while (low <= high) {
6:         int mid = (low + high) / 2;
7:         int midVal = a[mid];
8:
9:         if (midVal < key)
10:            low = mid + 1
11:         else if (midVal > key)
12:            high = mid - 1;
13:         else
14:            return mid; // key found
15:     }
16:     return -(low + 1); // key not found.
17: }
```

Fix:

$\text{low} + ((\text{high} - \text{low}) / 2)$

Vulnerable?

Read more:

<http://googleresearch.blogspot.co.uk/2006/06/extra-extra-read-all-about-it-nearly.html>

From the research dept

- <http://www.phrack.org/issues.html?issue=60&id=10>
- D. Brumley et al., [RICH: Automatically protecting against integer-based vulnerabilities](#), NDSS 2007
 - Formal typing rules for safe integer operations
 - Extend compiler to add dynamic checks when program violates type safety
- T. Wang et al., [IntScope: Automatically Detecting Integer Overflow Vulnerability in X86 Binary Using Symbolic Execution](#), NDSS 2009
 - Symbolic execute of binary code to find integer overflow
- D. Molnar et al., [Dynamic Test Generation To Find Integer Bugs in x86 Binary Linux Programs](#), USENIX Security 2009
 - SmartFuzz: tool to perform symbolic execution and dynamic test generation

Defenses

- Manual defenses
 - Avoid mixing signed/unsigned
 - Avoid mixing different width
 - Check for overflow/underflow
- SafeInt
 - Used extensively throughout Microsoft
 - <http://safeint.codeplex.com/>

ALGORITHMIC COMPLEXITY VULNERABILITIES

Denial of service attacks

- DoS attacks are based on the idea of overwhelming the target with the sheer number of requests
 - “smurf” ICMP attack
 - DNS amplification
- Other approach: are there special requests that cause the target to perform a disproportionate amount of work?
- Direction: look for requests that cause typically efficient algorithms to perform badly
 - Algorithmic complexity


Algorithmic complexity

- Data structures frequently used in applications have “average-case” expected running time that is far more efficient than worst case
 - They are used because typical/benign inputs do not expose the worse case
- Examples
 - Hash tables can degenerate to linked lists
 $O(1) \rightarrow O(n)$ insertion/lookup time of one element
($O(n^2)$ for n element)
 - Binary trees can degenerate to linked lists
 $O(\log n) \rightarrow O(n)$
- Can the attacker supplies inputs to force the algorithm into the worst case expected running time?

Hash tables


Normal

Buckets



Worst-case

Buckets



Hash tables

- Reduce object to a (say) 32-bit hash value
- Identify the bucket to store the object
 - Hash value modulo the bucket count
- If two input objects map to the same bucket, a *collision* has occurred
- Buckets hold linked lists of objects (whose hash value modulo the bucket count is the same)
 - Hash chains

Collisions

Two reasons for collisions

- Hash values are identical
 - Find k_1, k_2, \dots, k_i inputs such that $\text{Hash}(k_1) = \text{Hash}(k_2) = \dots = \text{Hash}(k_i)$
 - Hash collisions
- Hash values are different but their values modulo the bucket count are identical
 - Find k_1, k_2, \dots, k_i inputs such that $f(k_1) = f(k_2) = \dots = f(k_i)$ where f is the function mapping inputs to a bucket (e.g., $f(x) = \text{Hash}(x) \pmod{n}$)

Attacking hash functions

- Different hash functions
 - Cryptographic hash functions: MD5, SHA-1, etc.
 - Non cryptographic hash functions: functions with 32 bits of internal state (speed!), e.g., XOR
- Attacking weak (non cryptographic) hash functions
 - Weakness: limited internal state
 - Idea: find input (“*generator*”) such that internal state after hashing is the same as the initial state, say 0
 - Example: k_1, k_2, \dots, k_i such that
$$0 = \text{Hash}(k_1) = \text{Hash}(k_2) = \dots = \text{Hash}(k_i)$$
 - Then: any combination of k_1, k_2, \dots, k_i will be hashed to 0:
$$\text{Hash}(k_1 k_2) = \text{Hash}(k_2 k_1 k_3 k_2)$$
- Finding generators: bruteforce

Attack on Perl hash function

- Crosby and Wallach on Perl 5.6.1 and Perl 5.8.0
- Hash function: state machine with 32 bit state
 - Input is mixed in one byte at a time, using addition, multiplication, shifting
- Finding generators: 46 generators in 1 hour time
- Inputs: ~90k strings of 24 characters
- Loading time
 - Malicious inputs: 6,506 seconds (1:48 hours)
 - Benign inputs: 2 seconds

Attack on Java hashCode function

```
uint32_t hash(const char *arKey,
              uint32_t nKeyLength)
{
    uint32_t h = 0;

    for (; nKeyLength > 0;
          nKeyLength -=1)
    {
        h= ((h << 5) - h) + *arKey++;
    }

    return h;
}
```

- Equivalent substrings
 - hash('E^y') = 2260
hash('F^z') = 2260
 - hash('E^ya') = 70157
hash('F^za') = 70157
 - hash(abcE^ydef) = 2758600642447
hash(abcF^zdef) = 2758600642447
- Easy to generate infinite number of colliding strings (binary permutations)

Attacks

- Look at where hash tables are used:
 - IP addresses
 - Transaction IDs
 - Databases
 - URL parameters
 - HTTP headers
 - ...
- See A. Klink and J. Walde, [Effective Denial of service attacks against web application platforms](#), 28C3 ([MOV](#))

Solving algorithmic complexity attacks on hash tables

- Better hashing functions
 - Cryptographic hash functions
 - Universal hash functions
 - Guarantee that given k_1 and k_2 , the odds that $h(k_1) = h(k_2)$ are less than small value ϵ
 - Performance?
- Better data structures
 - Guaranteed runtime bounds, regardless of their inputs (e.g., red-black tree)
 - Implementation complexity?


Regular expressions

- Notation for describing set of characters
 - Literal characters
 - Metacharacters (*, +, ?)
- `abc*`
 - Matches `ab`, `abc`, `abcc`
 - Does not match `abcd`


Regular expressions and automata

- Matching $a(bb)^+a$

DFA



NFA



Examples courtesy of <http://swtch.com/~rsc/regexp/regexp1.html>

Matching regular expressions

- Backtracking
 - Whenever multiple paths are possible, take one (randomly) and record the choice
 - If matching fails, backtrack to last choice point (branch) and take alternative path
- The number of paths to check can be exponential in the length of the input string
- Attackers can create strings that force this worst-case backtracking
- Why using backtracking
 - Easy to implement back references

Attacks on regular expressions

- **(a|aa)*b**

| | |
|---|---------------|
| aac | 1.0209710598 |
| aac | 2.69697809219 |
| aac | 7.27811908722 |
| aac | 20.1901729107 |
| aac | 50.6117949486 |

- **[a-z]+@[a-z]+([a-z\.]+\.)+[a-z]+**

| | |
|--------------|----------------|
| redos@x..... | 0.192698001862 |
| redos@x..... | 0.820160150528 |
| redos@x..... | 3.50762486458 |
| redos@x..... | 5.67420506477 |
| redos@x..... | 9.44989609718 |
| redos@x..... | 15.4427139759 |
| redos@x..... | 41.7495079041 |

Detecting vulnerable regex

Statically

- Understand why a regular expression causes blowup in running time
- Search for regular expressions that have the same characteristics
- James Kirrage, Asiri Rathnayake, Hayo Thielecke, [Static Analysis for Regular Expression Denial-of-Service Attacks](http://arxiv.org/abs/1301.0849), 2013
<http://arxiv.org/abs/1301.0849>

Detecting vulnerable regex

Fuzzing:

- Generate inputs likely to cause issues with a given regex
- Measure the processing time
- Alert if it gets above a given threshold or increases faster than desired
- Microsoft, SDL Regex Fuzzer, <http://www.microsoft.com/en-us/download/details.aspx?id=20095>

Take-home points

- Denial of service attacks exploit imbalance between work done by attacker and work done by victim
- Computational complexity denial of service
 - Exploit algorithms that have bad worst-case running time...
 - By providing inputs that force the worst case to happen
- Examples: hash tables, regular expressions

Read more

- Solar Designer, [Designing and Attacking Port Scan Detection Tools](http://www.phrack.org/issues.html?issue=53&id=13#article), Phrack 1998
<http://www.phrack.org/issues.html?issue=53&id=13#article>
 - First mention of DoS attacks against hash tables
- S. Crosby and D. Wallach, [Denial of Service via Algorithmic Complexity Attacks](http://static.usenix.org/event/sec03/tech/full_papers/crosby/crosby_html/), USENIX Security 2003
http://static.usenix.org/event/sec03/tech/full_papers/crosby/crosby_html/
 - First thorough discussion; found vulnerabilities in Perl, Bro (IDS), Squid
- A. Klink and J. Walde, [Efficient Denial of Service Attacks on Web Application Platforms](http://events.ccc.de/congress/2011/Fahrplan/attachments/2007_28C3_Effective_DoS_on_web_application_platforms.pdf), 28C3 2011
http://events.ccc.de/congress/2011/Fahrplan/attachments/2007_28C3_Effective_DoS_on_web_application_platforms.pdf
 - Break PHP, Python, Java, Ruby
- J. Aumasson et al., [Hash-flooding DoS reloaded: attacks and defenses](https://131002.net/siphash/siphashdos_appsec12_slides.pdf), APPSEC 2012
https://131002.net/siphash/siphashdos_appsec12_slides.pdf
<http://emboss.github.com/blog/2012/12/14/breaking-murmur-hash-flooding-dos-reloaded/>
 - More sophisticated analysis of hash functions (differential cryptoanalysis)
- P. Junod, [Like a Hot Knife Through Butter](http://crypto.junod.info/2012/12/13/hash-dos-and-btrfs/), 2012
<http://crypto.junod.info/2012/12/13/hash-dos-and-btrfs/>
 - Attacking the btrfs filesystem

Read more

- OWASP, https://www.owasp.org/index.php/Regular_expression_Denial_of_Service_-_ReDoS
 - Examples of real-world, problematic regular expressions
- R. Smith et al., [Backtracking Algorithmic Complexity Attacks Against a NIDS](http://www.acsac.org/2006/papers/54.pdf), ACSAC 2006
<http://www.acsac.org/2006/papers/54.pdf>
 - Attack against Snort intrusion detection system
- K. Namjoshi and G. Narlikar, Robust and Fast Pattern Matching For Intrusion Detection, INFOCOM 2010
<http://ect.bell-labs.com/who/knamjoshi/papers/robustness-infocom10.pdf>
 - More problems in Snort (back references)

Next time

- Sample exam